

CS 4501: Introduction to Computer Vision

Filtering and Edge Detection

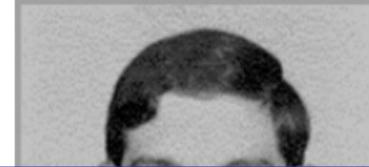
Connelly Barnes

Outline

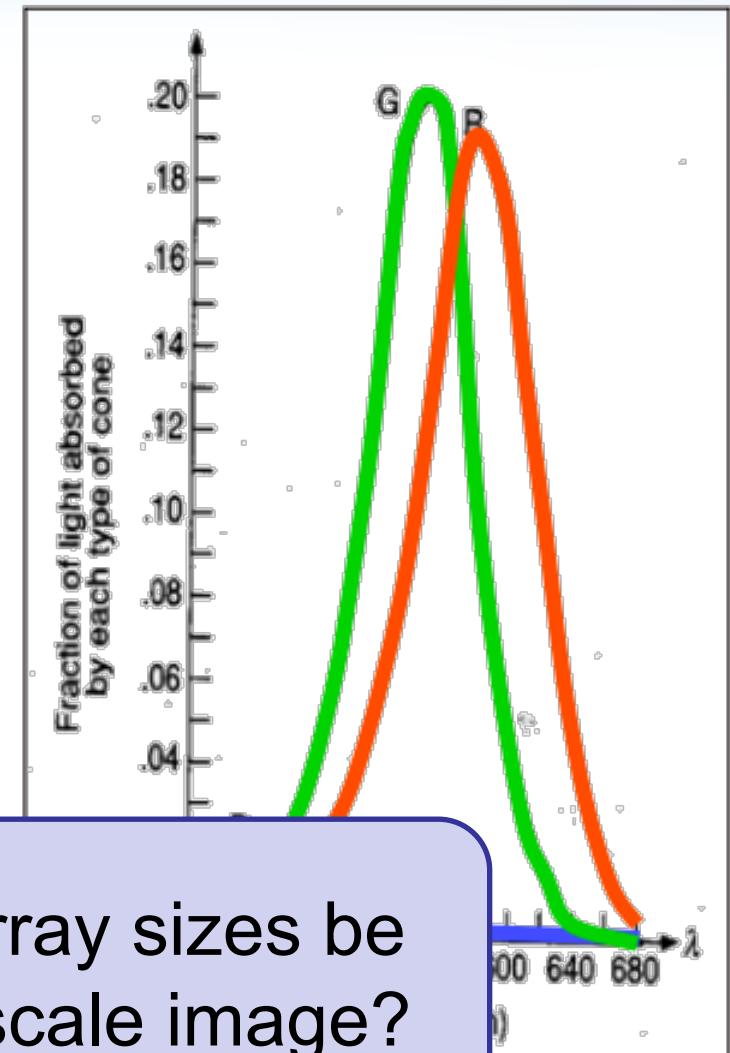
- Simple image processing
 - Greyscale
 - Brightness
 - Mirroring or “flipping”
- Filtering
 - Linear filters: cross-correlation and convolution
 - Gaussian filters
- Edge detection
 - Simple edge detector
 - Canny edge detector

Image processing: greyscale

- The human retina perceives red, green, blue.
- To compute luminance of a pixel, we need to take an average of the RGBs:
 - $L = 0.21 R + 0.72 G + 0.07 B$ (ITU HDTV)
 - $L = 0.3 R + 0.59 G + 0.11 B$ ([W3C](#))



If represented as arrays, what would the array sizes be for the input RGB image? The output greyscale image?



Original

Grayscale

Figure taken from

FvDFH

Image processing: brightness

- Simply scale the array of RGB values.
 - Must clamp to valid range [0, 1]
 - $I_{\text{out}} = \min(\alpha I_{\text{in}}, 1)$
- Where is this operation used?
 - Photo adjustment, dataset augmentation



Original



Brighter

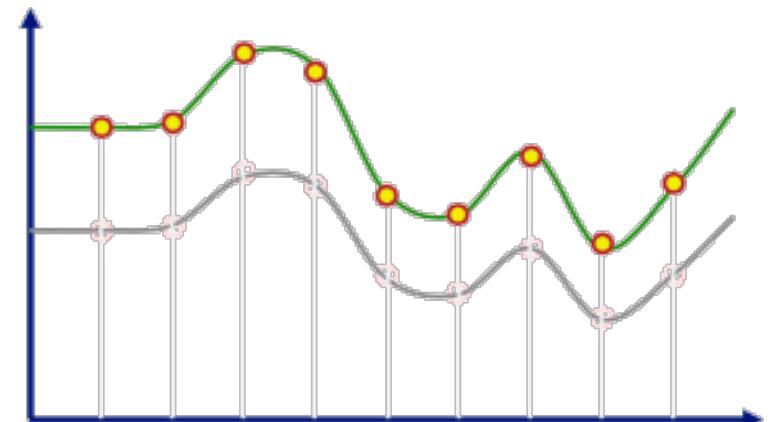


Image processing: mirroring or “flipping”



Why do we care?

Some linear filters involve “flipping” operations.
Another way to do dataset augmentation.

Vertically
flipped

Vertically
and
horizontally
flipped

Outline

- Simple image processing
 - Greyscale
 - Brightness
 - Mirroring or “flipping”
- **Filtering**
 - Linear filters: cross-correlation and convolution
 - Gaussian filters
- Edge detection
 - Simple edge detector
 - Canny edge detector

Image filtering

- **Filtering:**
 - Form a new image whose pixels are a combination of original pixel values.
- **Goals:**
 - Extract useful information from image
 - Features (corners, edges, blobs, ...)
 - Enhance image properties
 - Remove noise, remove unwanted objects, ...

Image filtering

De-noising



Salt and pepper noise

Super-resolution



In-painting



Bertamio et al

Outline

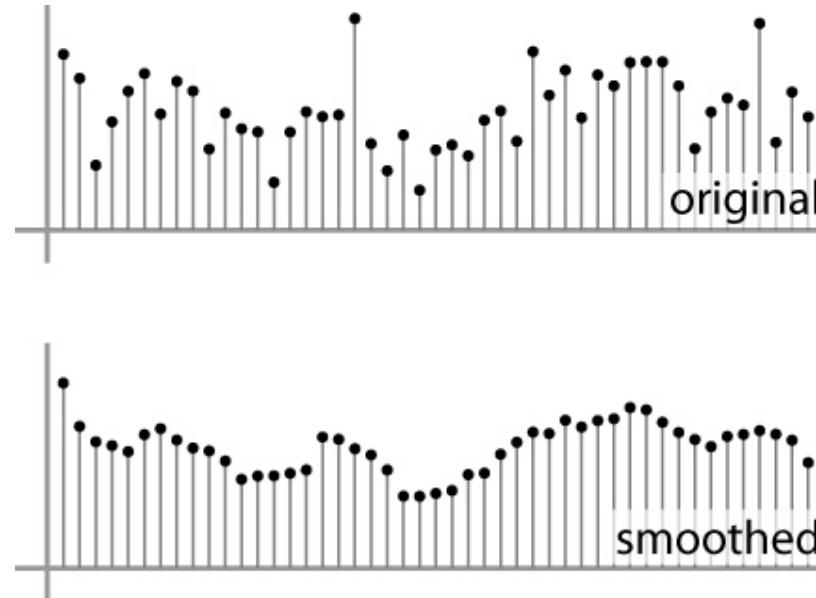
- Simple image processing
 - Greyscale
 - Brightness
 - Mirroring or “flipping”
- Filtering
 - **Linear filters: cross-correlation and convolution**
 - Gaussian filters
- Edge detection
 - Simple edge detector
 - Canny edge detector

Linear filtering: a key idea

- Transformations on signals; e.g.:
 - bass/treble controls on stereo
 - blurring/sharpening operations in image editing
 - smoothing/noise reduction in tracking
- Key properties
 - linearity: $\text{filter}(f + g) = \text{filter}(f) + \text{filter}(g)$
 - shift invariance: behavior invariant to shifting the input
 - delaying an audio signal
 - sliding an image around
- Can be modeled mathematically by *convolution*

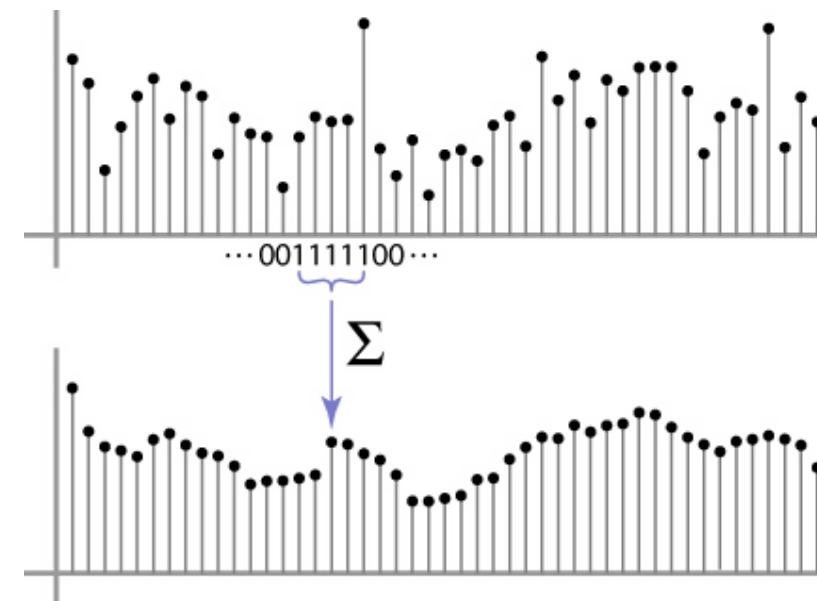
Moving Average

- basic idea: define a new function by averaging over a sliding window
- a simple example to start off: smoothing



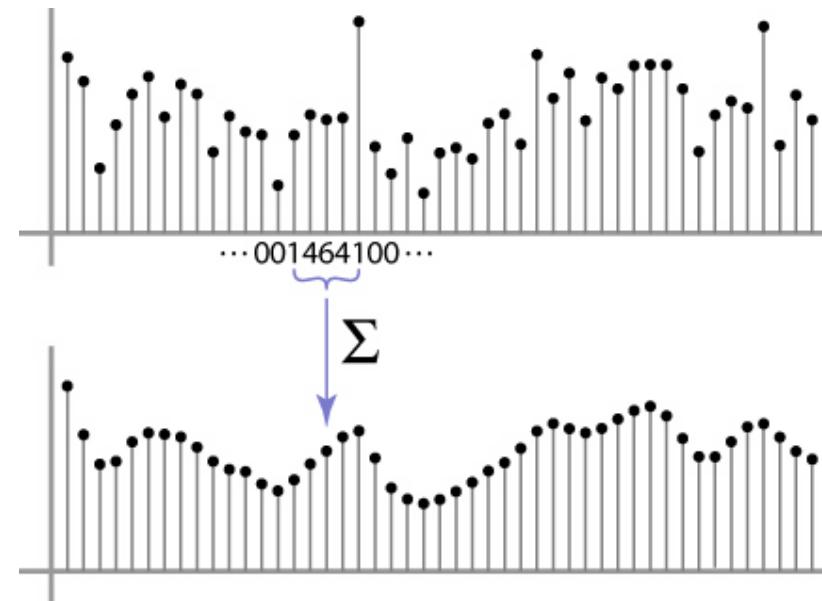
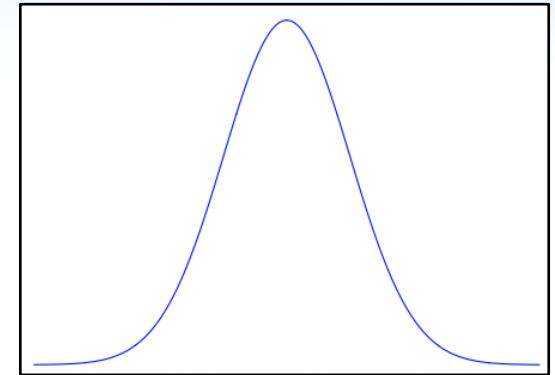
Weighted Moving Average

- Can add weights to our moving average
- Weights $[..., 0, 1, 1, 1, 1, 1, 0, ...] / 5$



Weighted Moving Average

- Bell curve (gaussian-like) weights $[..., 1, 4, 6, 4, 1, ...]$



Moving Average In 2D

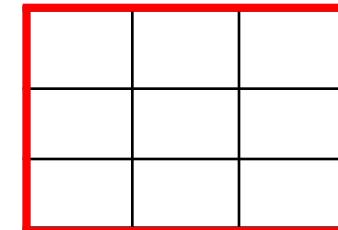
What are the weights H?

For uniform filter? (takes the mean)

For bell curve shaped filter?

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

Input image $F[x, y]$



$H[u, v]$

Cross-correlation filtering

- Let's write this down as an equation. Assume the averaging window is $(2k+1) \times (2k+1)$:

$$G[i, j] = \frac{1}{(2k+1)^2} \sum_{u=-k}^k \sum_{v=-k}^k F[i+u, j+v]$$

- We can generalize this idea by allowing different weights for different neighboring pixels:

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i+u, j+v]$$

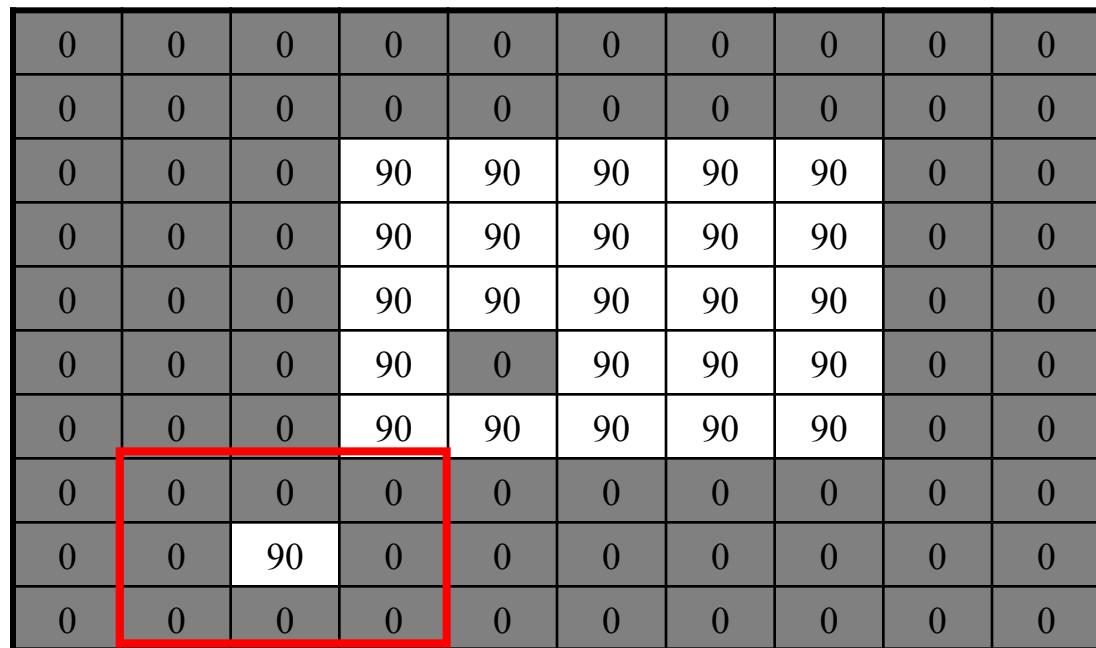
- This is called a **cross-correlation** operation and written:

$$G = H \otimes F$$

- H is called the “filter” or “kernel.”

Gaussian filtering

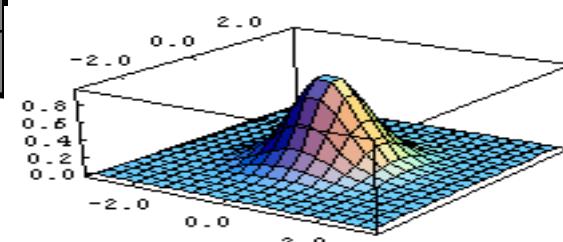
- A Gaussian kernel gives less weight to pixels further from the center of the window



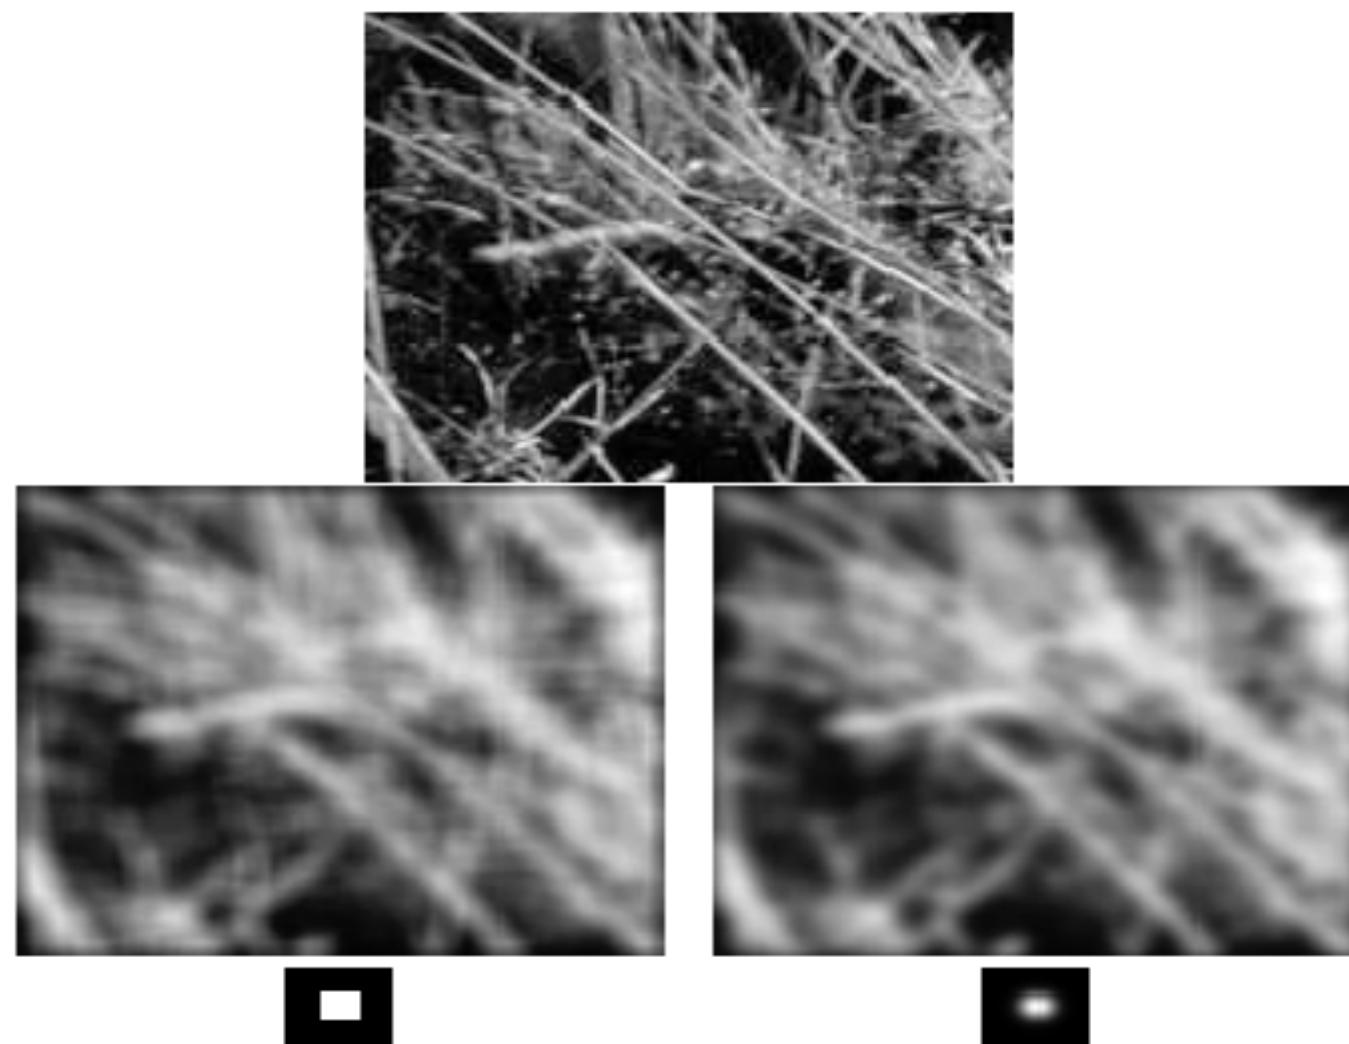
$$F[x, y]$$

$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$$

$$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$
$$H[u, v]$$



Box Filter vs. Gaussian Filter



Convolution

- **Cross-correlation:**

$$G = H \otimes F$$

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v]F[i + u, j + v]$$

- **Convolution** is similar to cross-correlation, but the filter is flipped horizontally and vertically before being applied:

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v]F[i - u, j - v]$$

- It is written:

$$G = H * F$$

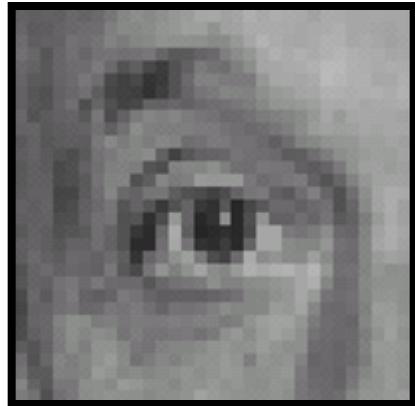
- Suppose H is a Gaussian or uniform (mean) kernel. How does convolution differ from cross-correlation?

Convolution is nice!

- Notation: $b = c * a$
 - Convolution is a multiplication-like operation
 - commutative
 - associative
 - distributes over addition
 - scalars factor out
 - identity: unit impulse $e = [..., 0, 0, 1, 0, 0, ...]$
 - Conceptually no distinction between filter and signal
 - Usefulness of associativity
 - often apply several filters one after another: $((a * b_1) * b_2) * b_3$
 - this is equivalent to applying one filter: $a * (b_1 * b_2 * b_3)$
- $$a * b = b * a$$
- $$a * (b * c) = (a * b) * c$$
- $$a * (b + c) = a * b + a * c$$
- $$\alpha a * b = a * (\alpha b) = \alpha(a * b)$$
- $$a * e = a$$

Practice with linear filters

- Assume we are using **cross-correlation** filtering (filter is **not** flipped)

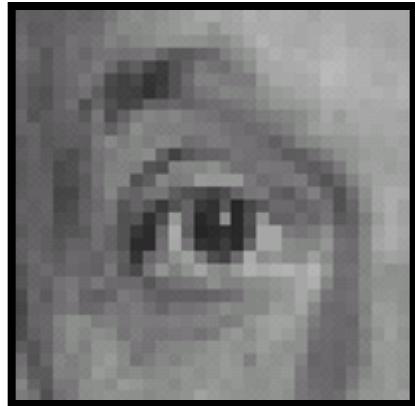


Original

0	0	0
0	1	0
0	0	0

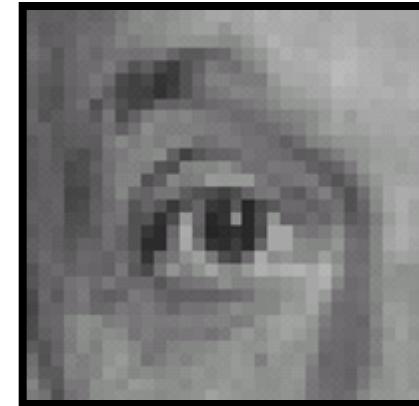
?

Practice with linear filters



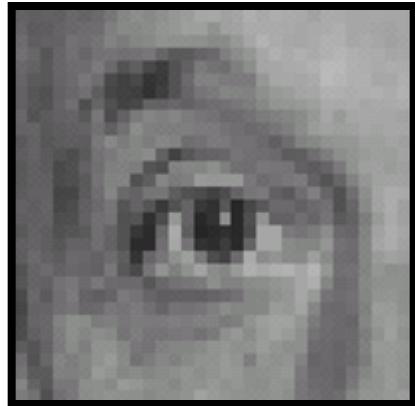
Original

0	0	0
0	1	0
0	0	0



Filtered
(no change)

Practice with linear filters

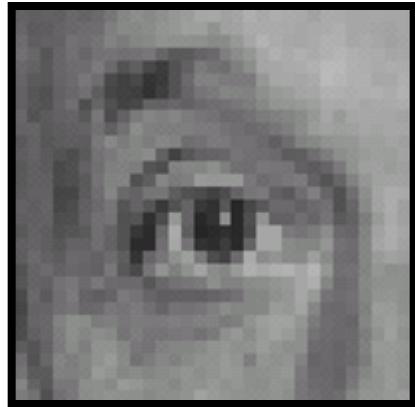


Original

0	0	0
0	0	1
0	0	0

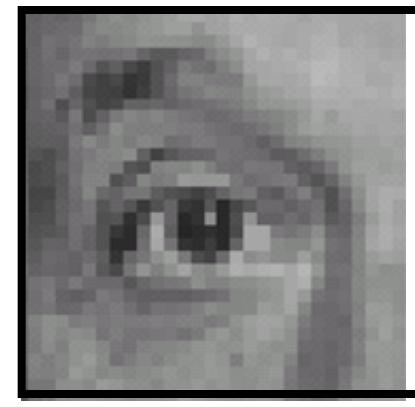
?

Practice with linear filters



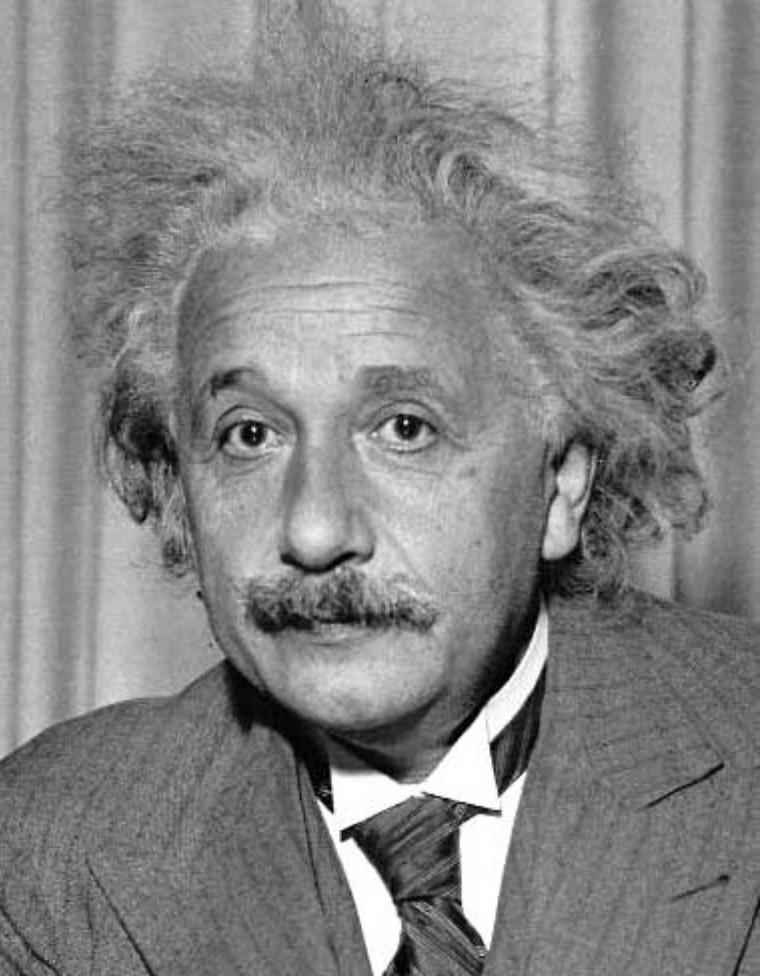
Original

0	0	0
0	0	1
0	0	0



Shifted left
By 1 pixel

Practice with linear filters

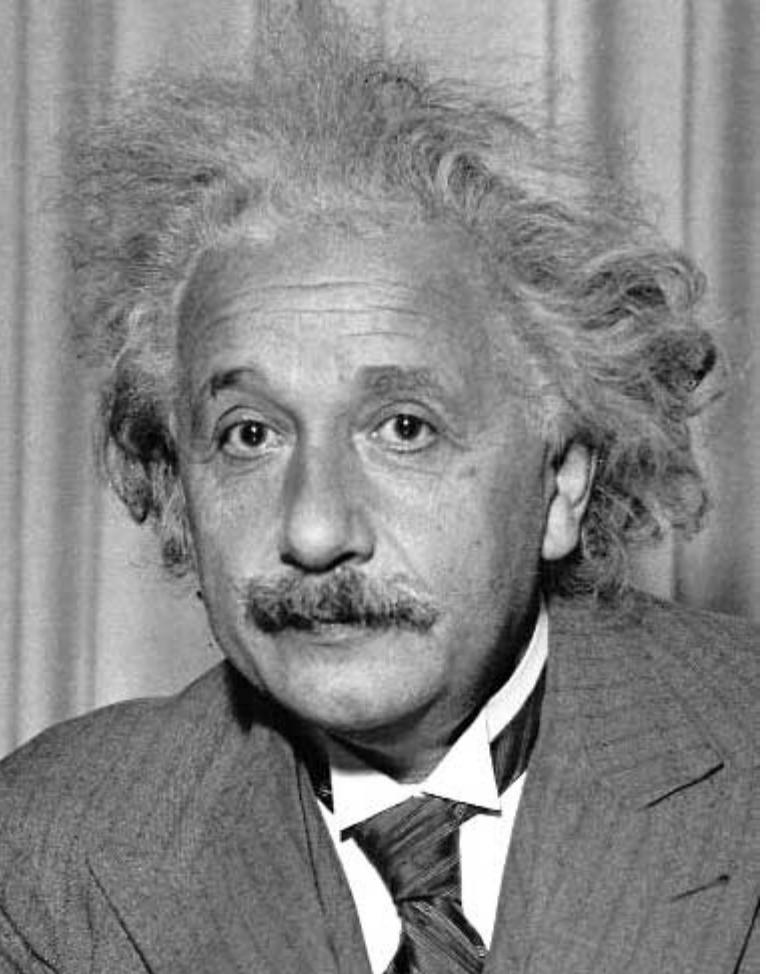


1	0	-1
2	0	-2
1	0	-1

Sobel

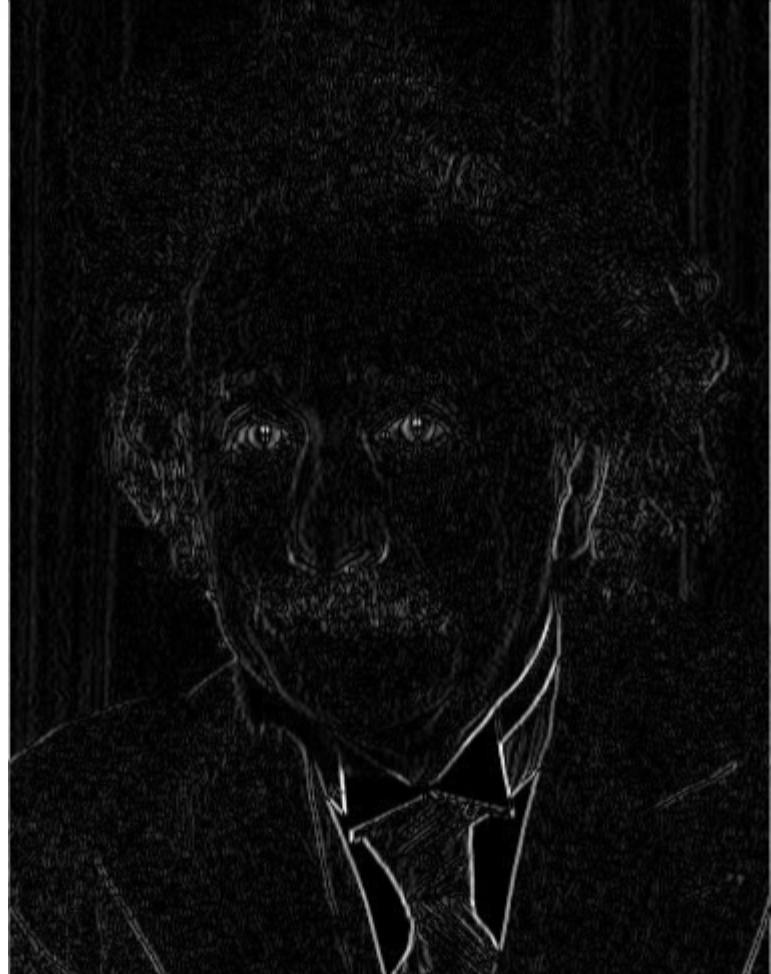
?

Practice with linear filters



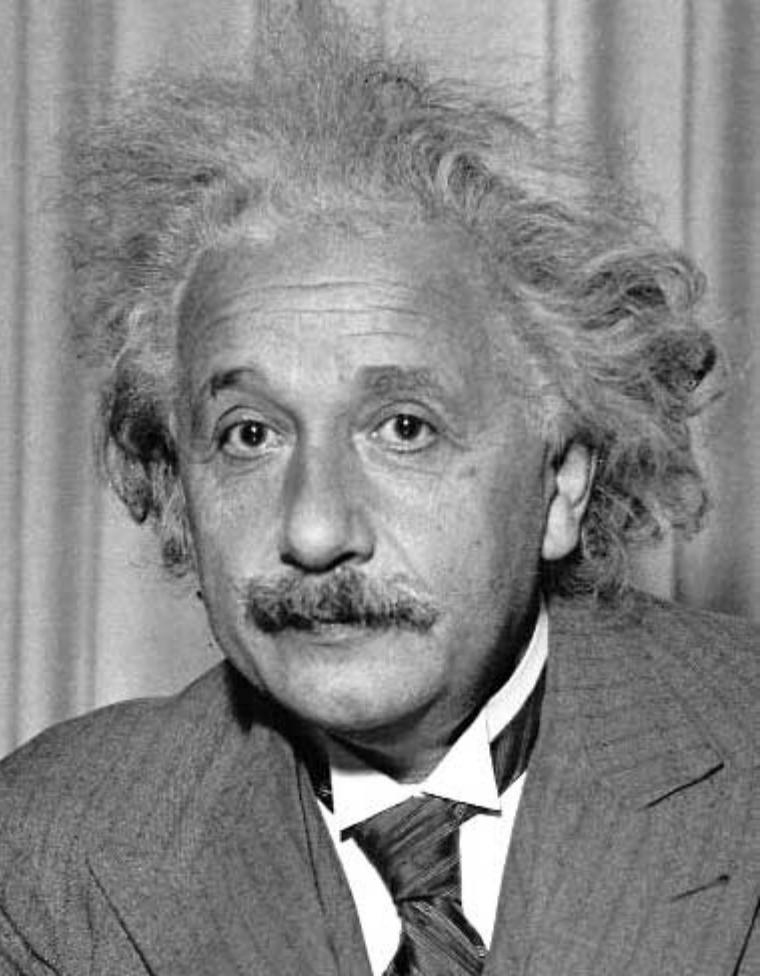
1	0	-1
2	0	-2
1	0	-1

Sobel



Vertical Edge
(absolute value)

Practice with linear filters

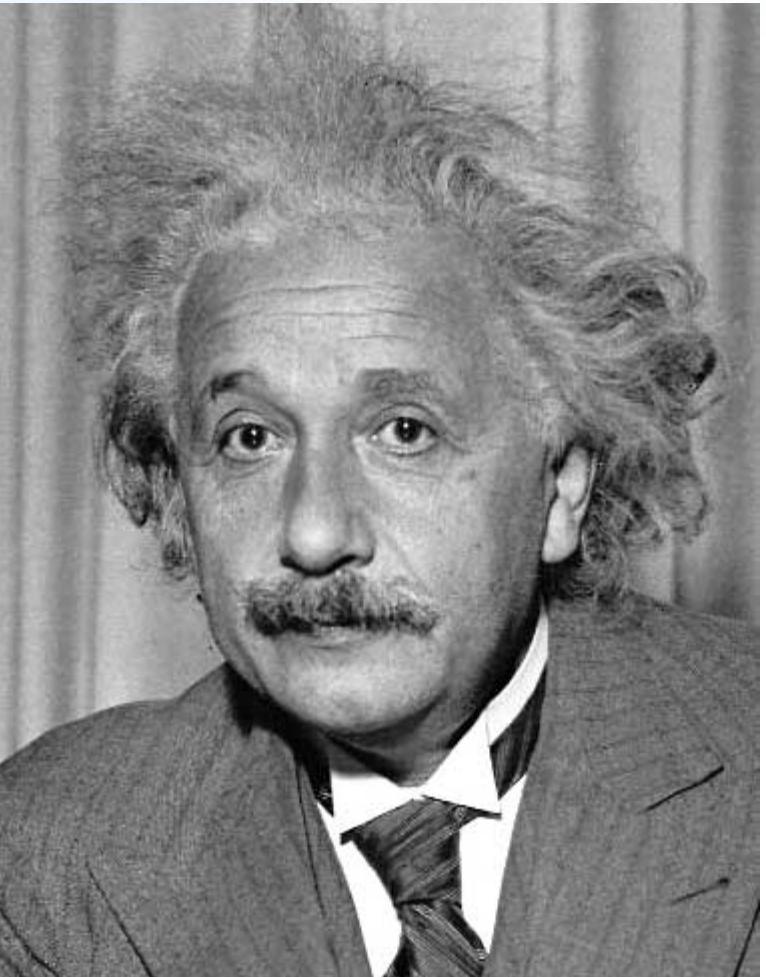


1	2	1
0	0	0
-1	-2	-1

Sobel

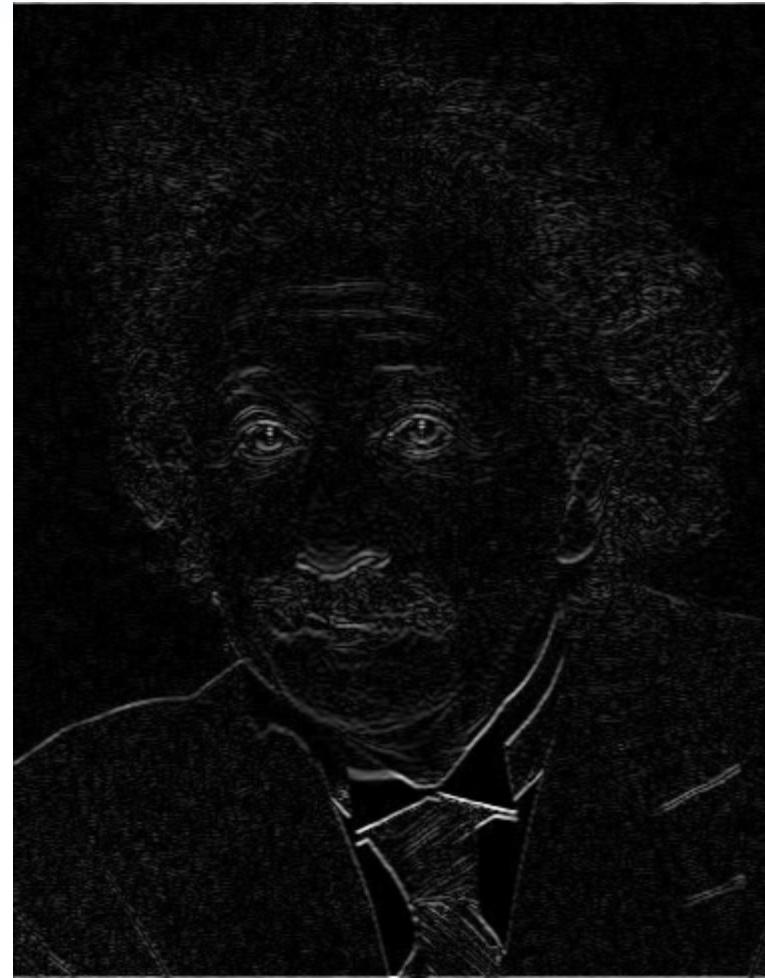
?

Practice with linear filters



1	2	1
0	0	0
-1	-2	-1

Sobel



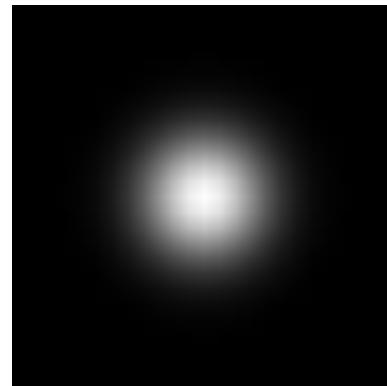
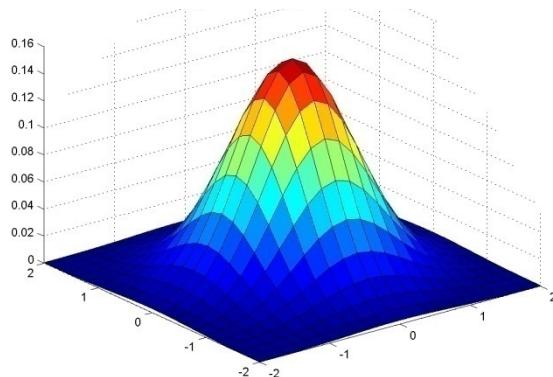
Horizontal Edge
(absolute value)

Outline

- Simple image processing
 - Greyscale
 - Brightness
 - Mirroring or “flipping”
- Filtering
 - Linear filters: cross-correlation and convolution
 - **Gaussian filters**
- Edge detection
 - Simple edge detector
 - Canny edge detector

Important linear filter: Gaussian

- Weight contributions of neighboring pixels by nearness



0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

$5 \times 5, \sigma = 1$

$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

- Same shape in spatial and frequency domain
(Fourier transform of Gaussian is Gaussian)

Gaussian filters

- Remove “high-frequency” components from the image (low-pass filter)
 - Images become more smooth
- Convolution with self is another Gaussian
 - So can smooth with small-width kernel, repeat, and get same result as larger-width kernel would have
 - Convolving two times with Gaussian kernel of width σ is same as convolving once with kernel of width $\sigma\sqrt{2}$

Gaussian filters



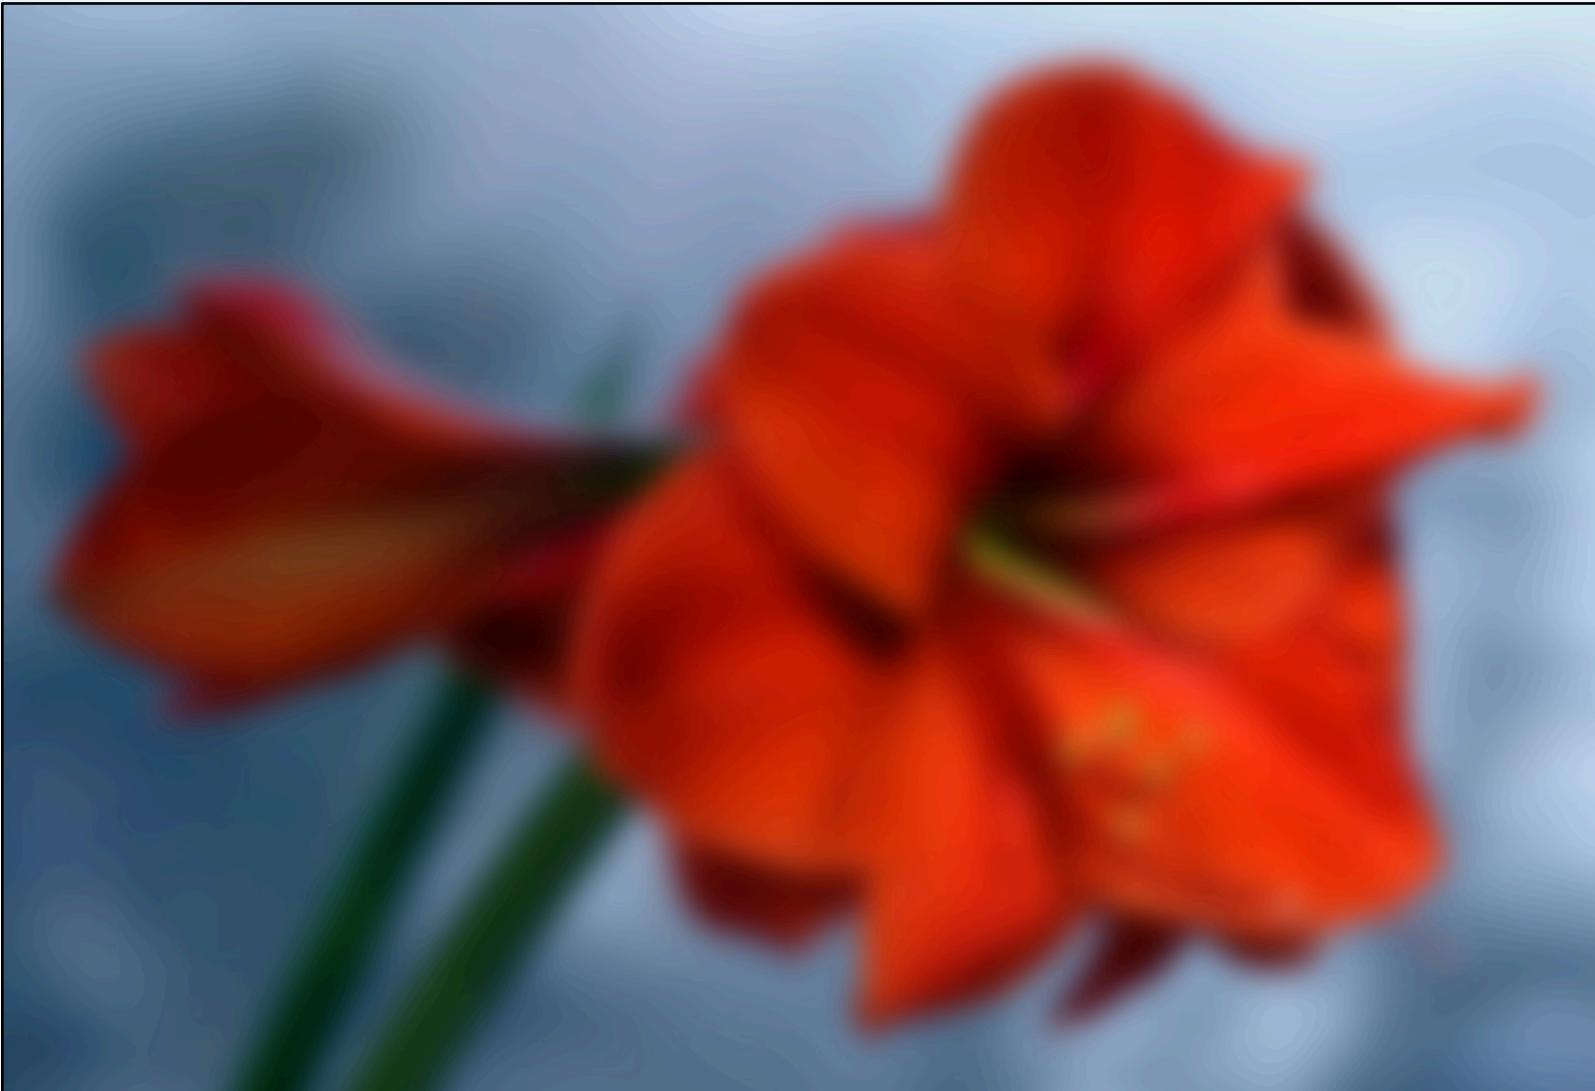
Input image (2048 x 1397)

Gaussian filters



Gaussian filtered ($\sigma=5$)

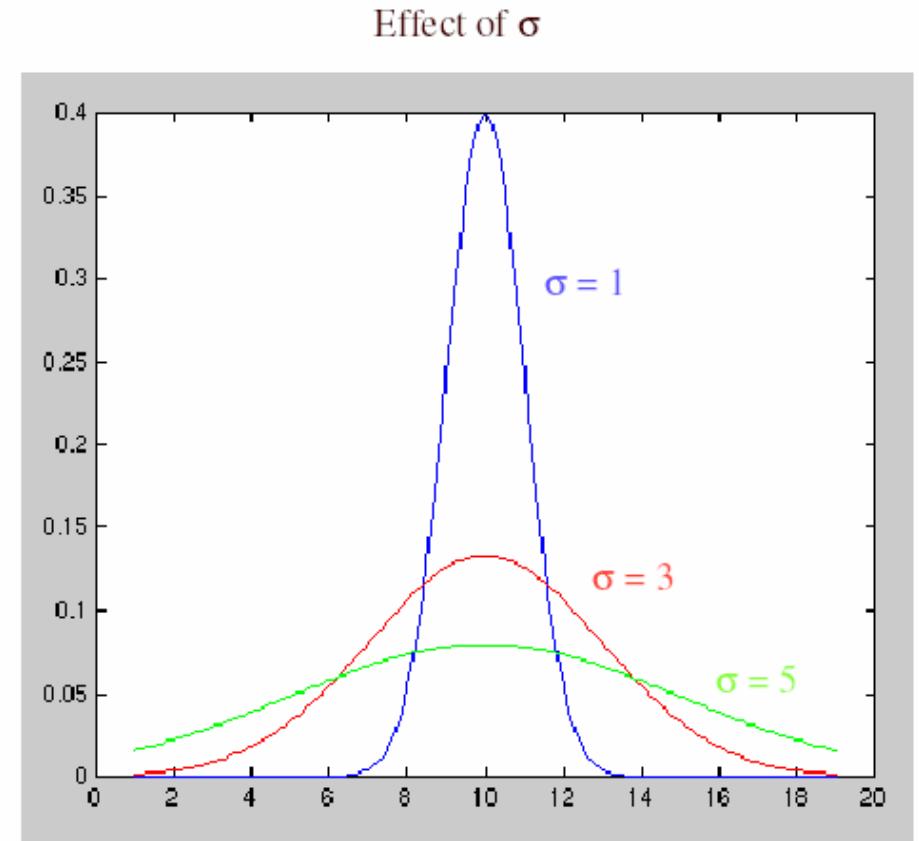
Gaussian filters



Gaussian filtered ($\sigma=20$)

Practical matters

- How big should the filter be?
 - Values at edges should be near zero
 - Rule of thumb for Gaussian: set filter half-width to about 3σ
 - Normalize truncated kernel. Why?



Separable Filters

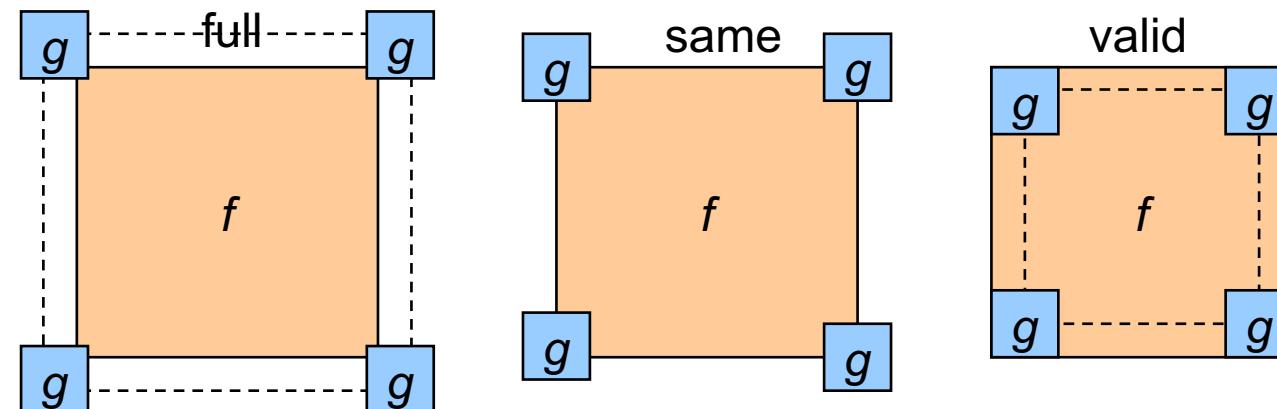
- Some kernels K can be written:

$$K = H * V, \quad H \text{ is horizontal, } V \text{ is vertical}$$

- Example: 2D Gaussian
$$\begin{aligned} G(x, y) &= \frac{1}{Z} \exp\left[\frac{-(x^2 + y^2)}{2\sigma^2}\right] \\ &= \frac{1}{Z} \exp\left[\frac{-x^2}{2\sigma^2}\right] \exp\left[\frac{-y^2}{2\sigma^2}\right] = \frac{1}{Z} H(x)V(y) \end{aligned}$$
- Filter first by H then V (or vice versa)
- Why is this useful?

Size of Output

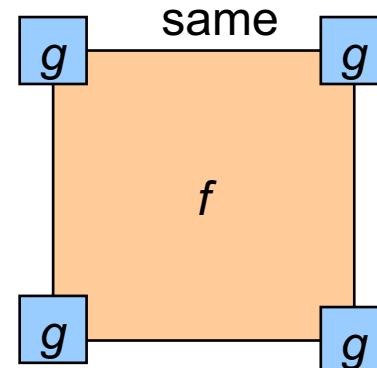
- MATLAB: `conv2(g,f,shape)`
- Python: `scipy.signal.convolve2d(g,f,shape)`
 - *shape = ‘full’*: output size is sum of sizes of f and g
 - *shape = ‘same’*: output size is same as f
 - *shape = ‘valid’*: output size is difference of sizes of f, g
- Easier for color images: `scipy.ndimage.filters.convolve(g,f)`



Source: S. Lazebnik

Python convolution (with SciPy)

- Python: `scipy.signal.convolve2d(g,f,shape)`
 - Convolves 2D images (e.g. greyscale)
- Python: `scipy.ndimage.filters.convolve(g,f)`
 - Convolves n-D images (e.g. greyscale, color)
 - But always uses ‘same’ size output
 - Can specify how to handle out of bounds pixels (e.g. ‘constant’, ‘reflect’)



Demo in Python

- Python (Jupyter) Notebook

Python Environment for Programming Assignments

- Recommend Python (another option: MATLAB?)
- Set up Python (recommend [Anaconda Python](#))
- Already included packages: SciPy, matplotlib, scikit-image.
- Recommend:
 - tensorflow (neural networks), ideally configured with GPU support
 - keras (neural networks)
- Can sign up to use department machines also,
will send document for that.

Outline

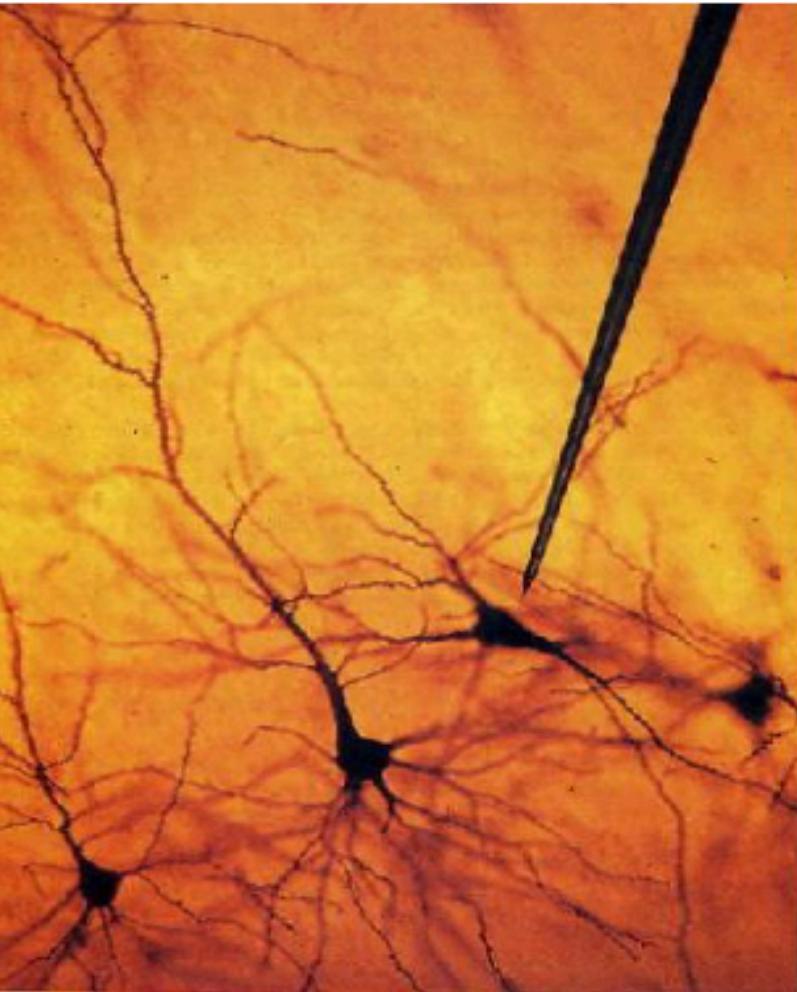
- Simple image processing
 - Greyscale
 - Brightness
 - Mirroring or “flipping”
- Filtering
 - Linear filters: cross-correlation and convolution
 - Gaussian filters
- **Edge detection**
 - Simple edge detector
 - Canny edge detector

Edge Detection

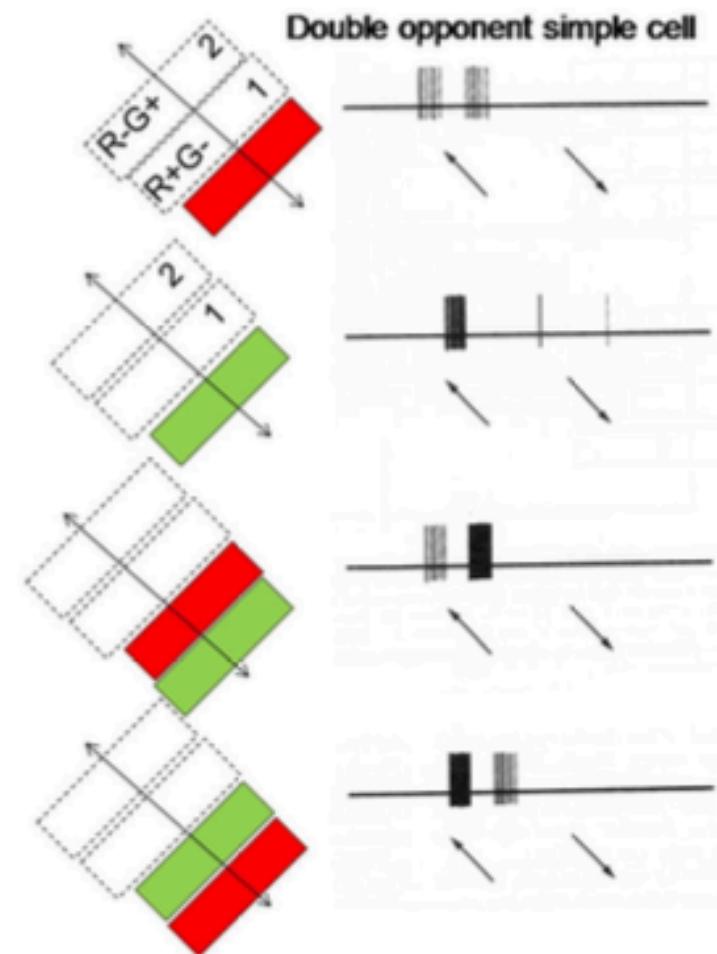


- (A) Cave painting at Chauvet, France, about 30,000 B.C.;
- (B) Aerial photograph of the picture of a monkey as part of the Nazca Lines geoglyphs, Peru, about 700 – 200 B.C.;
- (C) Shen Zhou (1427-1509 A.D.): Poet on a mountain top, ink on paper, China;
- (D) Line drawing by 7-year old I. Lleras (2010 A.D.).

Edge Detection: Mammal Vision



Hubel & Wiesel, 1960s



Edge Detection: Human Vision

152 Biederman

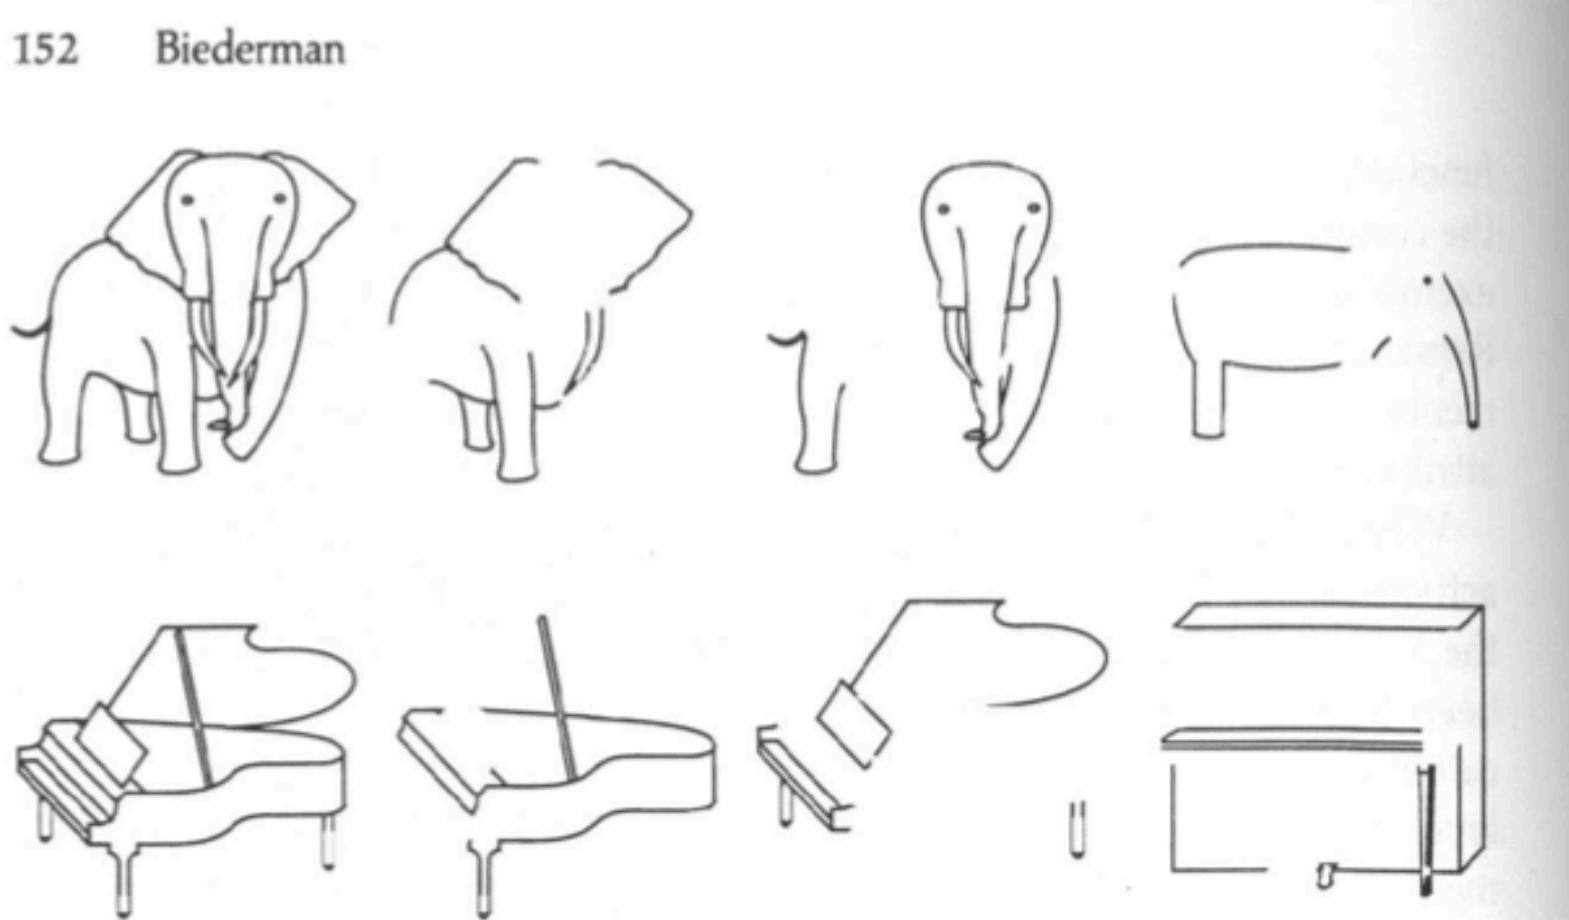
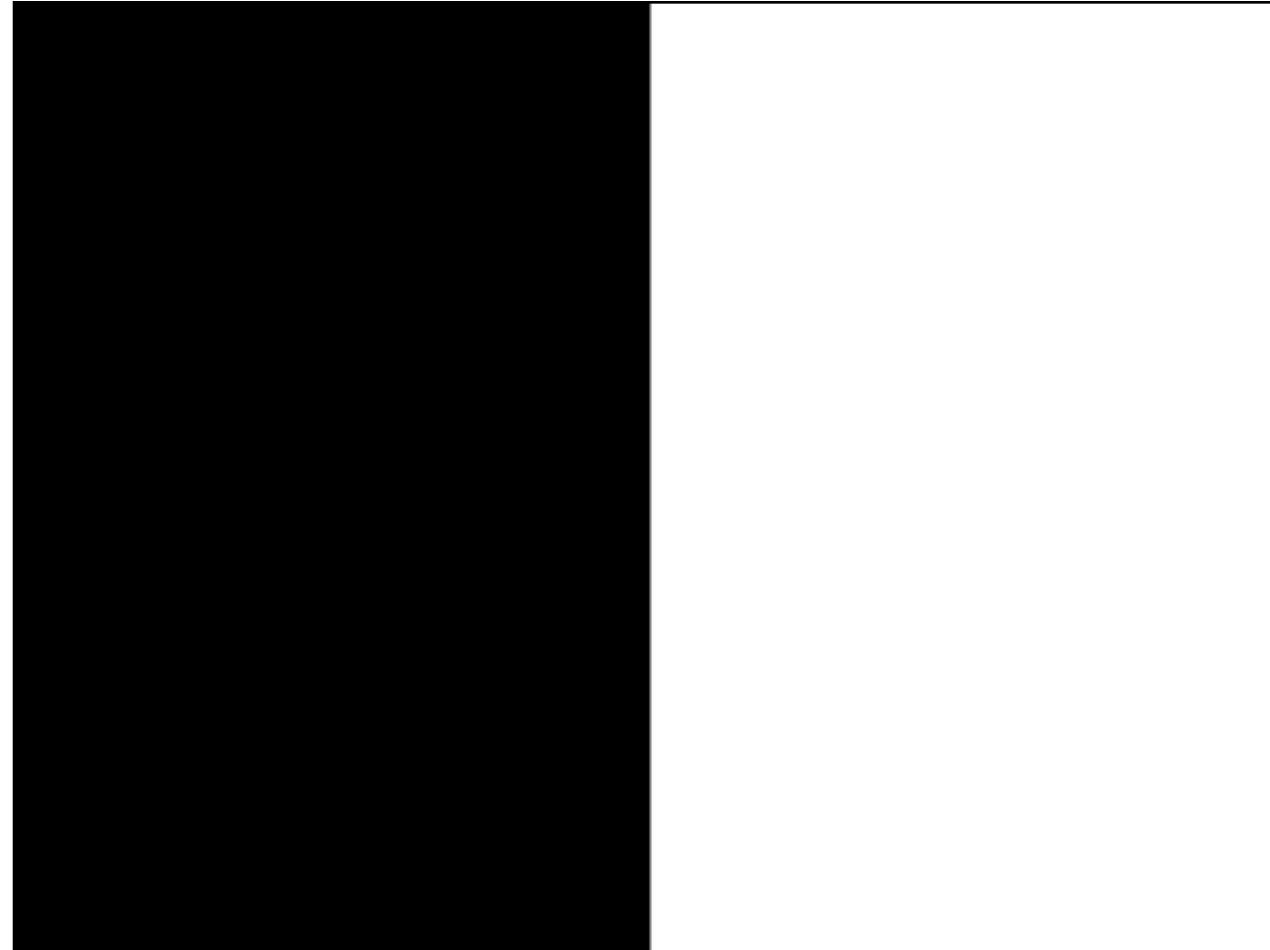


Figure 4.14

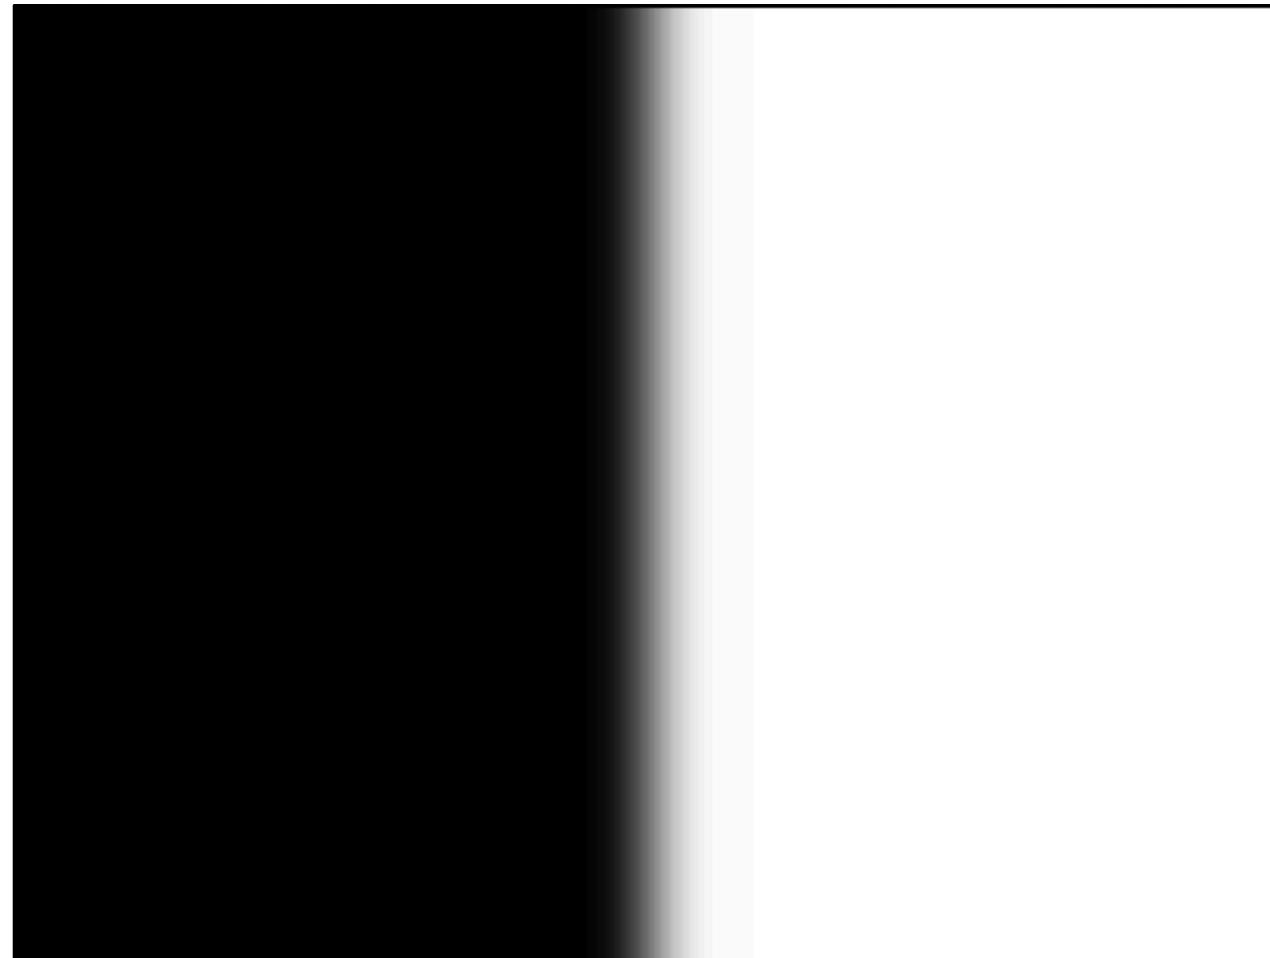
Complementary-part images. From an original intact image (left column), two complemen-

Slide from Fei Fei Li, Juan Carlos Niebles

What is an Edge?



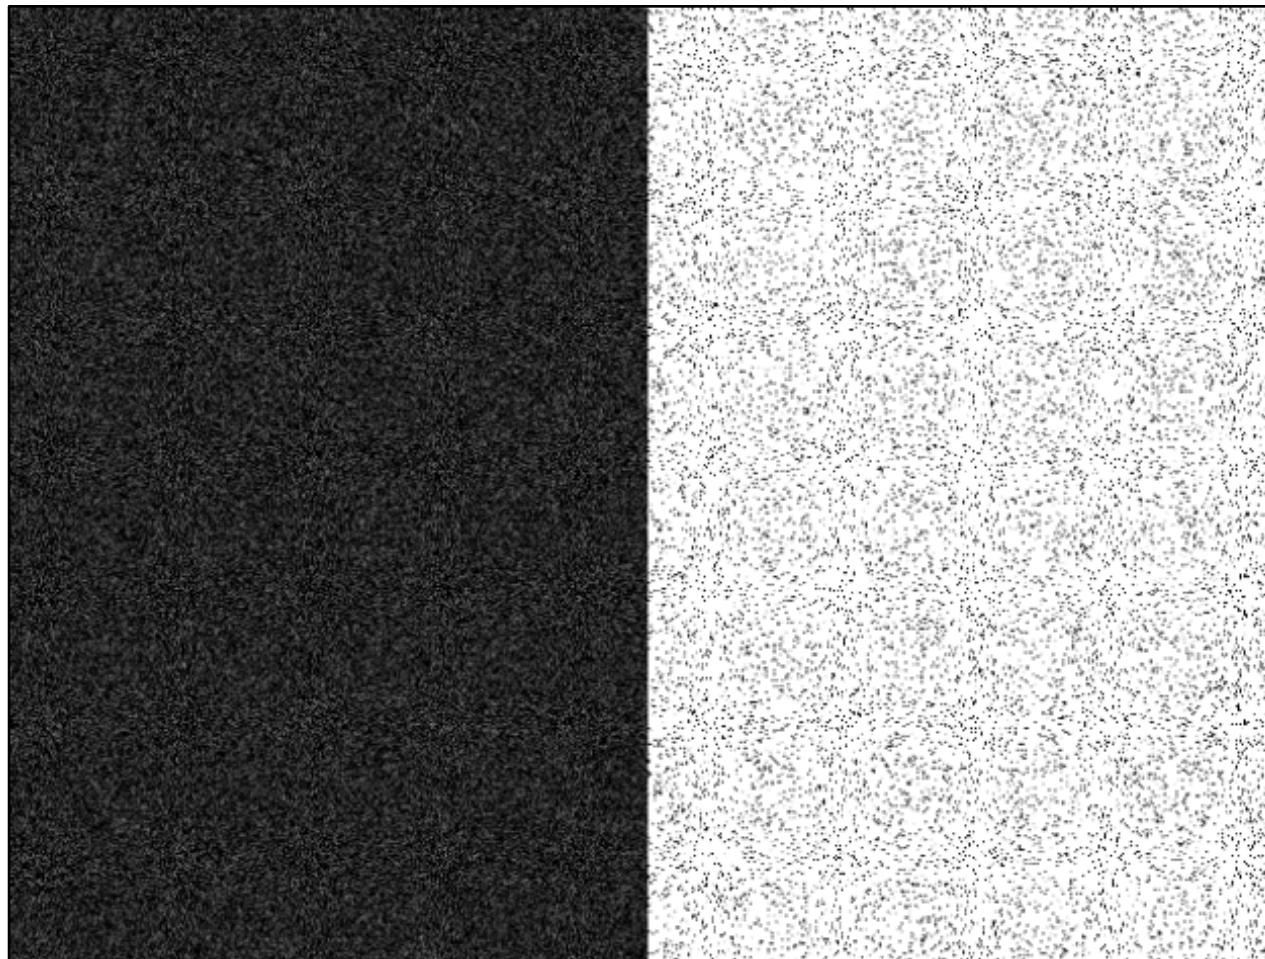
What is an Edge?



Challenge: blur

Slide from Jason Lawrence

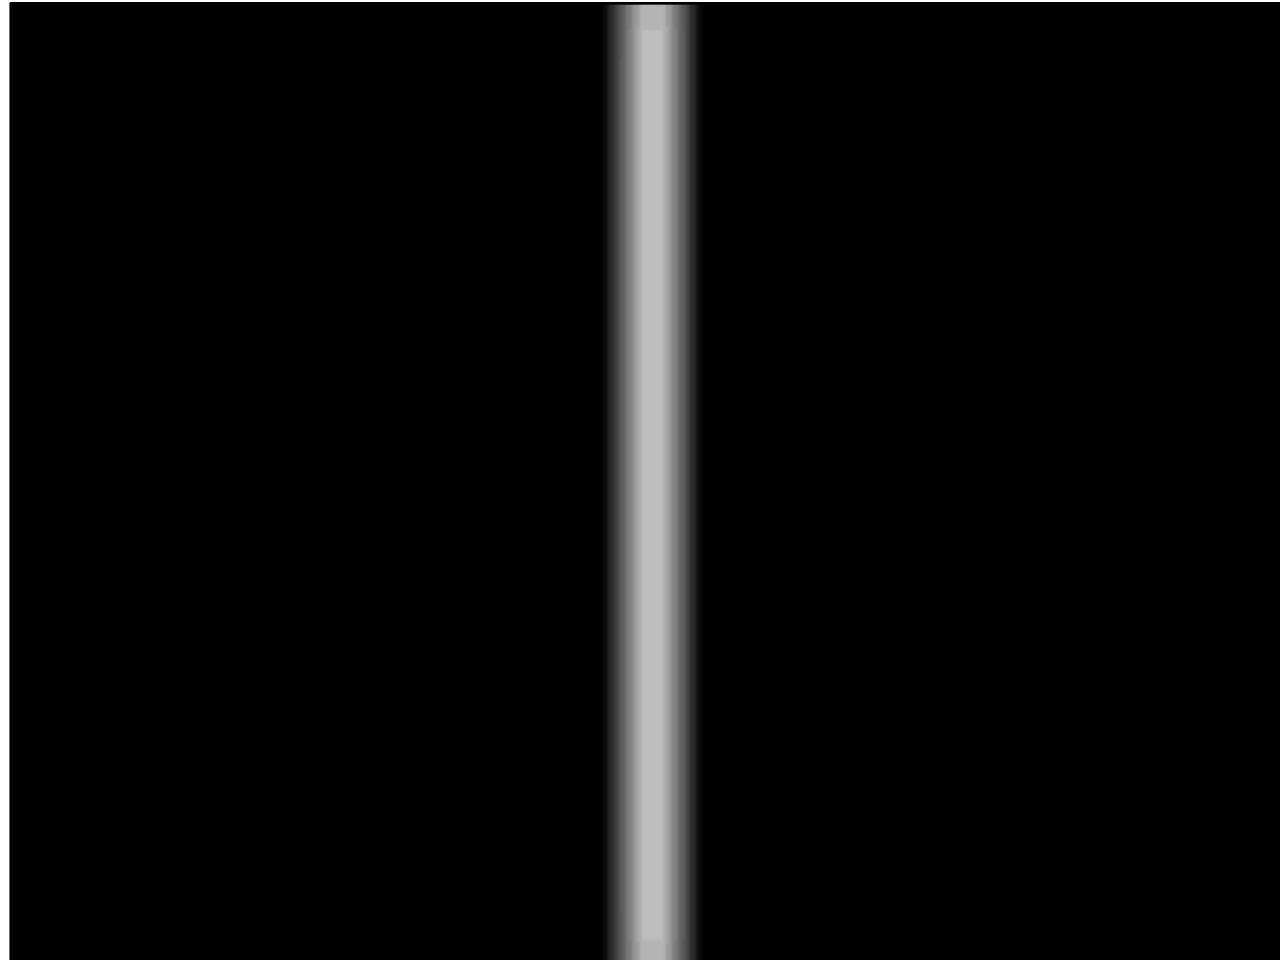
What is an Edge?



Challenge: noise

Slide from Jason Lawrence

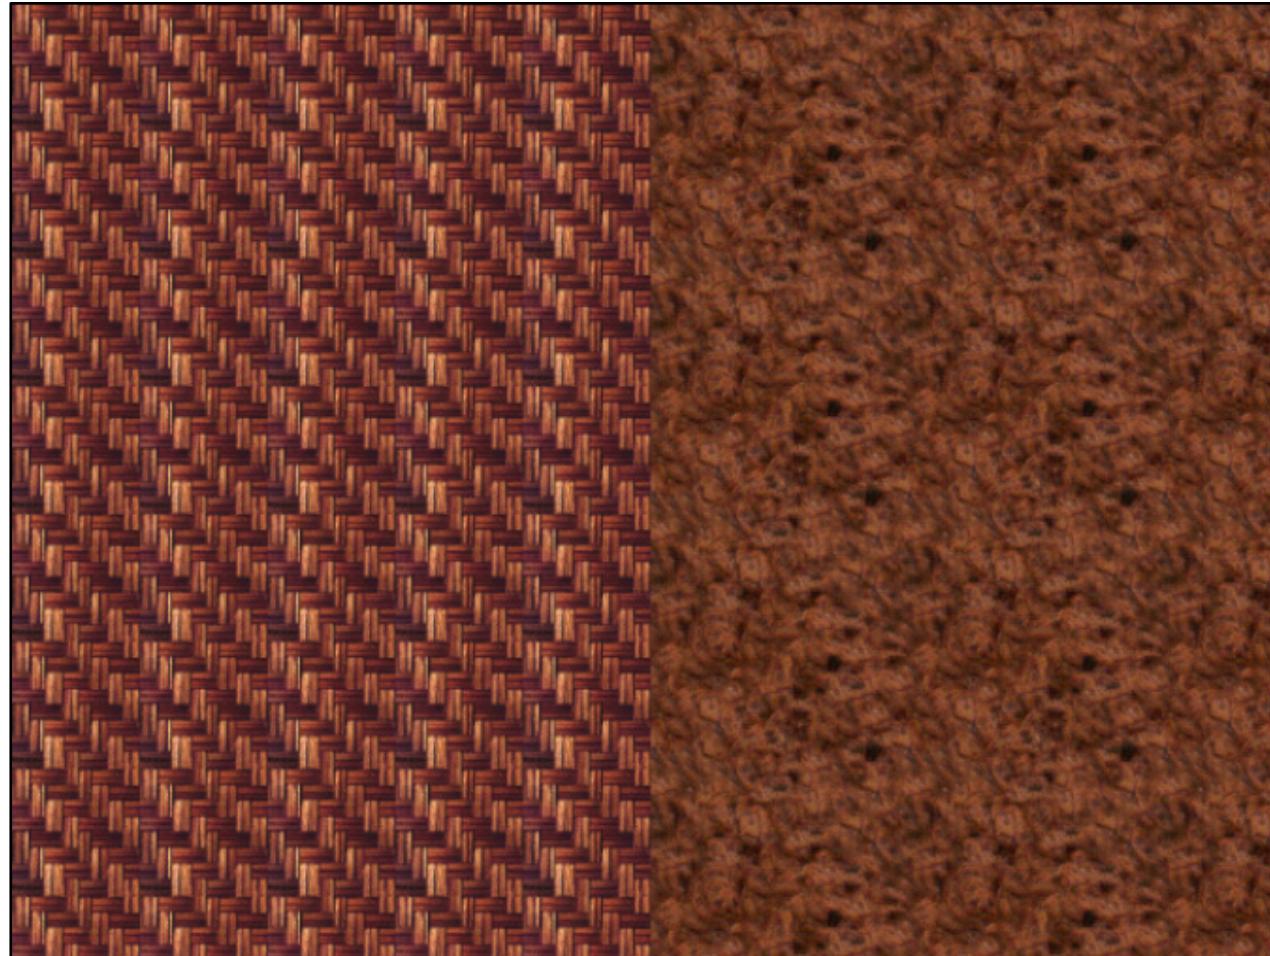
What is an Edge?



Is this one edge or two?

Slide from Jason Lawrence

What is an Edge?



Where are the edges?

Slide from Jason Lawrence

Characterizing Edges

- An **edge** is a place of rapid change in the image intensity function.

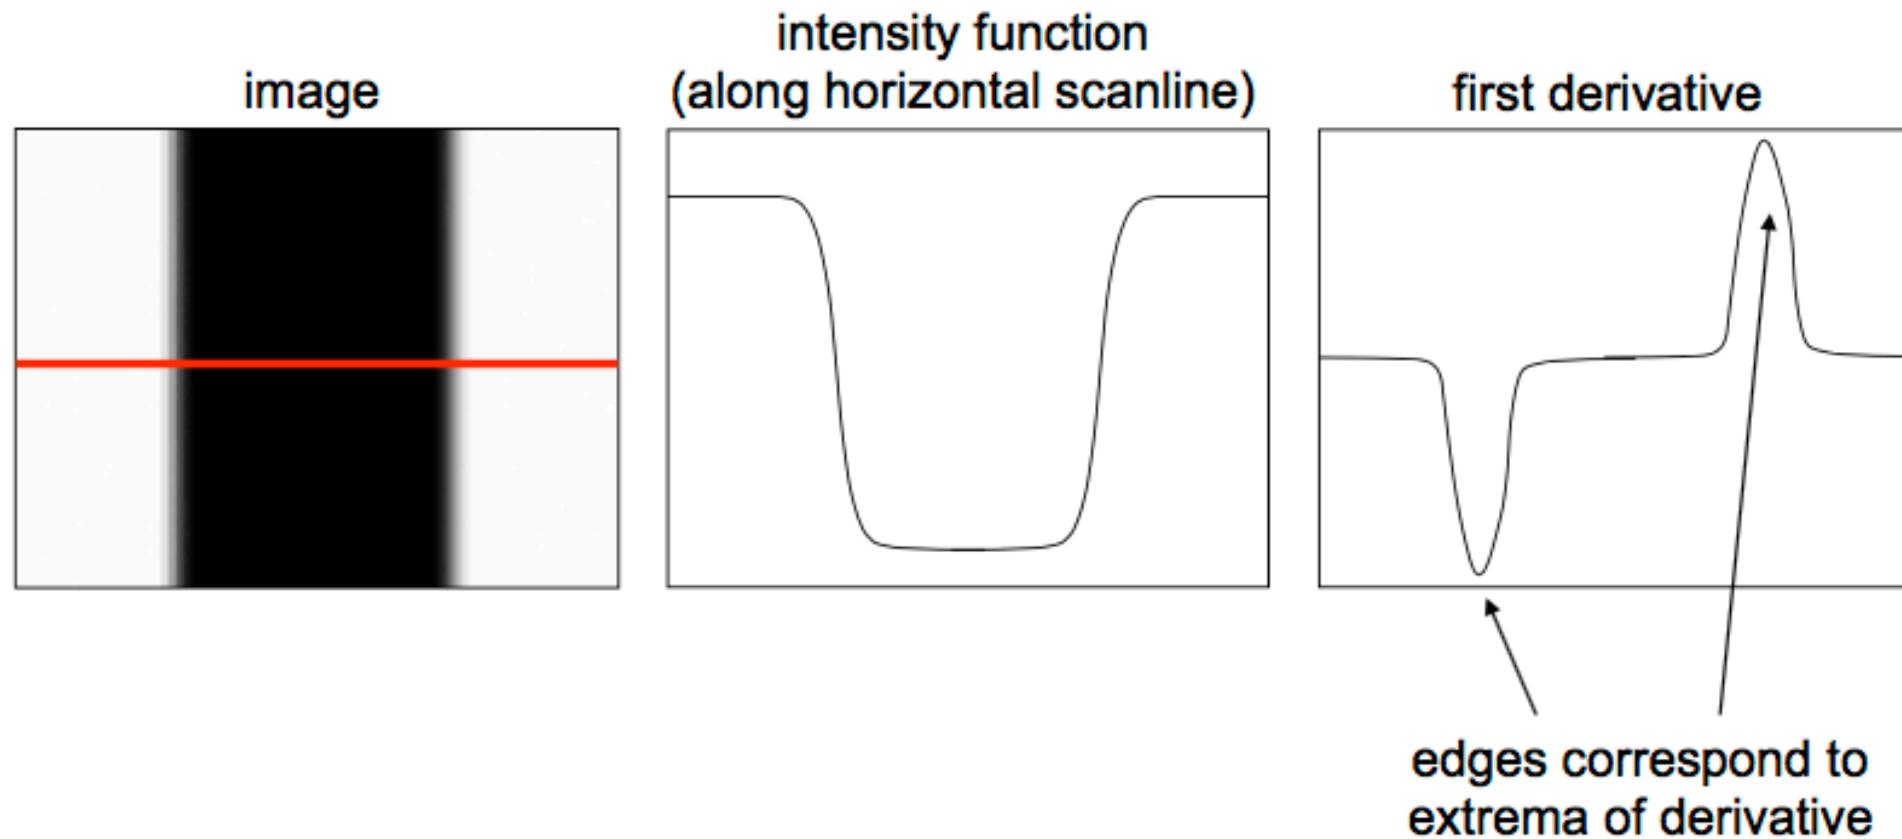
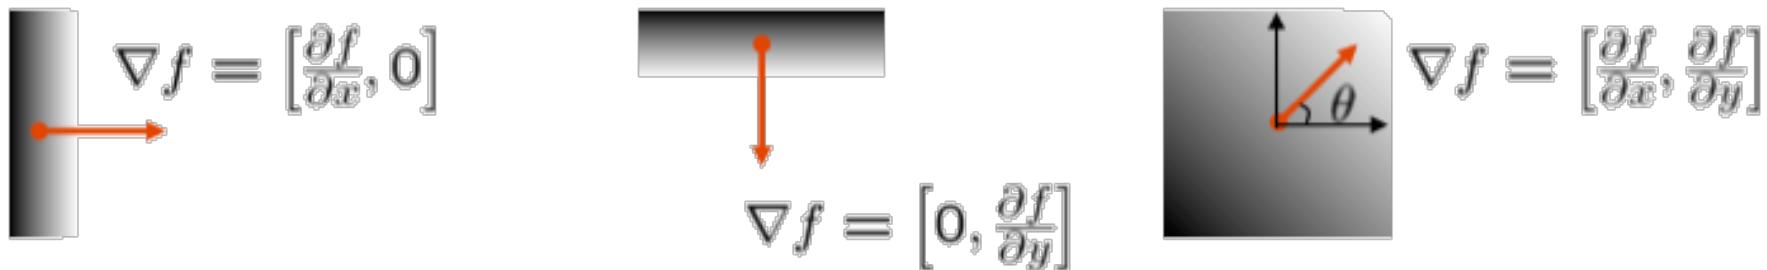


Image Gradient

- The gradient of an image: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$



The gradient vector points in the direction of most rapid increase in intensity

The gradient direction is given by $\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$

- how does this relate to the direction of the edge?

The *edge strength* is given by the gradient magnitude

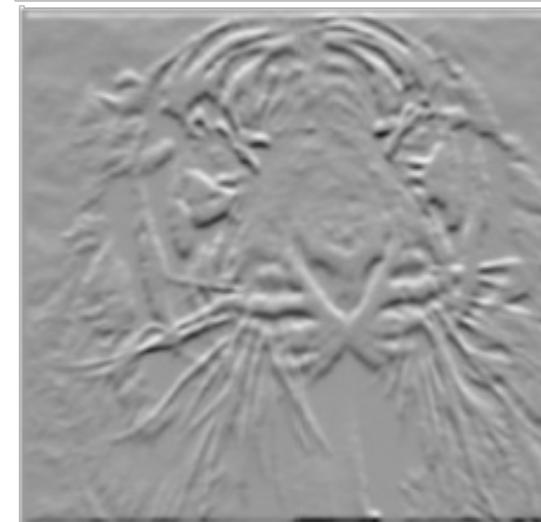
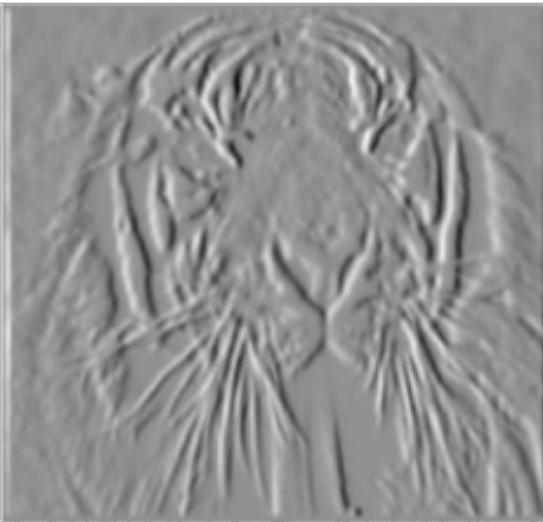
$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

Image Gradient

Original
Image



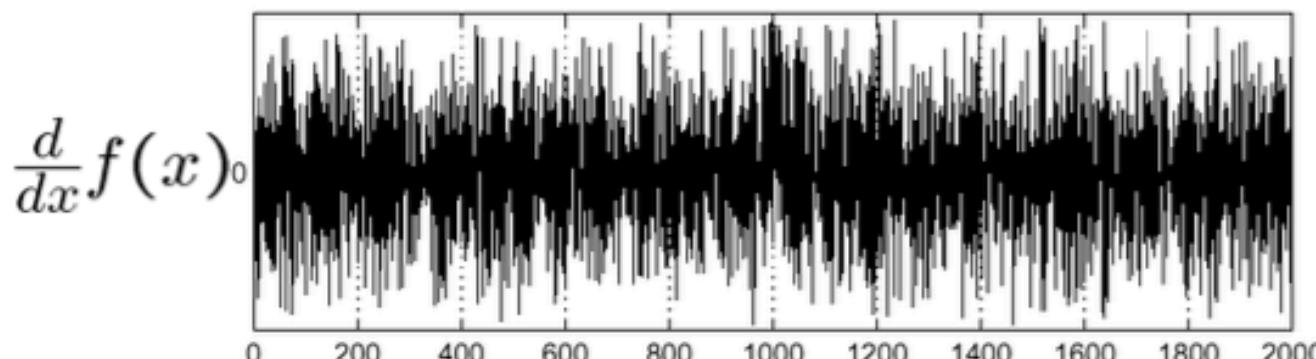
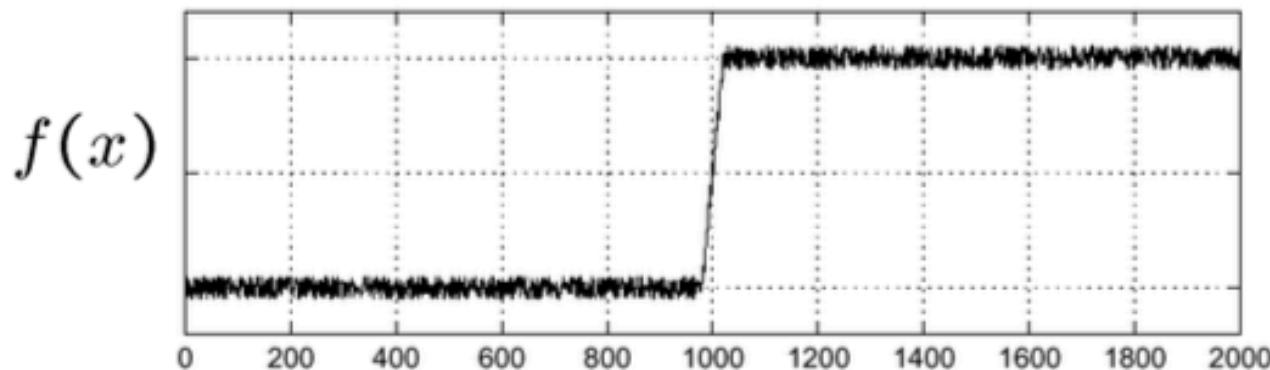
Gradient
magnitude



- Which one is the gradient in the x-direction? How about y-direction?

Effects of Noise

- Consider a single row or column of the image
 - Plotting intensity as a function of position gives a signal



Where is the edge?

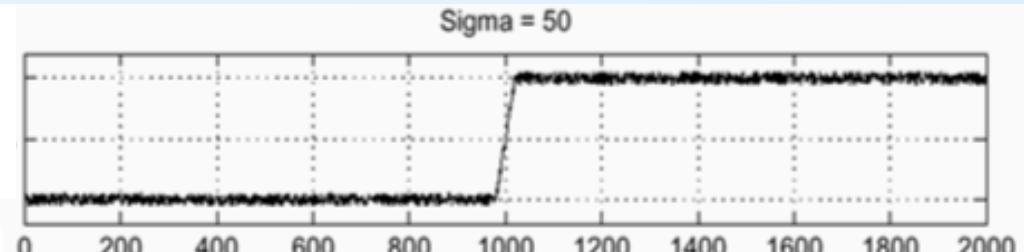
Source: S. Seitz

Simple Edge Detector

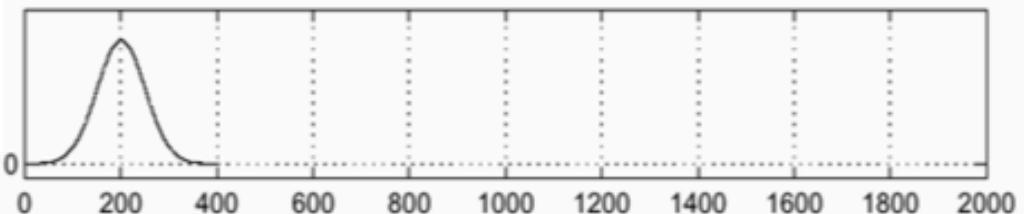
Algorithm:

1. Blur using Gaussian filter
2. Find gradient magnitude

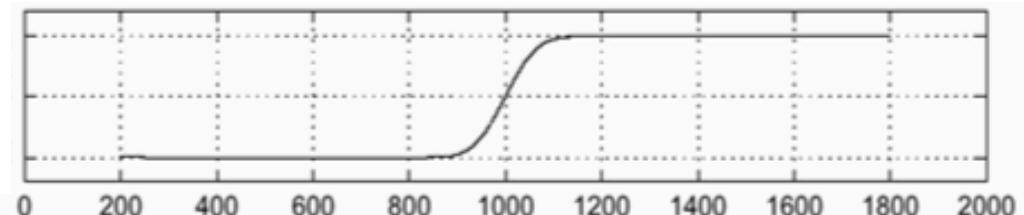
Input image: f



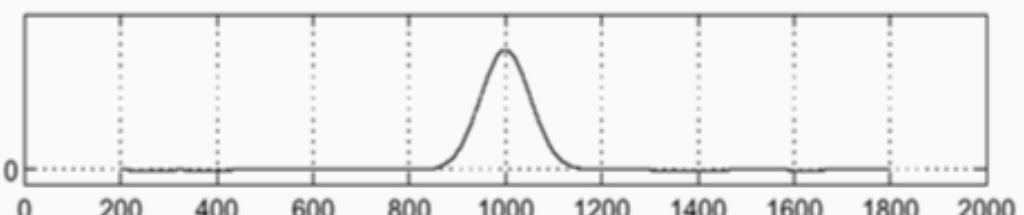
Gaussian kernel: g



Blurred image: $f * g$



Gradient: $\frac{d}{dx}(f * g)$



- To find edges, look for peaks in $\frac{d}{dx}(f * g)$

Derivatives of Filters

- Can optionally combine the blurring and differentiation steps using the theorem:

$$\frac{d}{dx}(f * g) = f * \frac{d}{dx}g$$

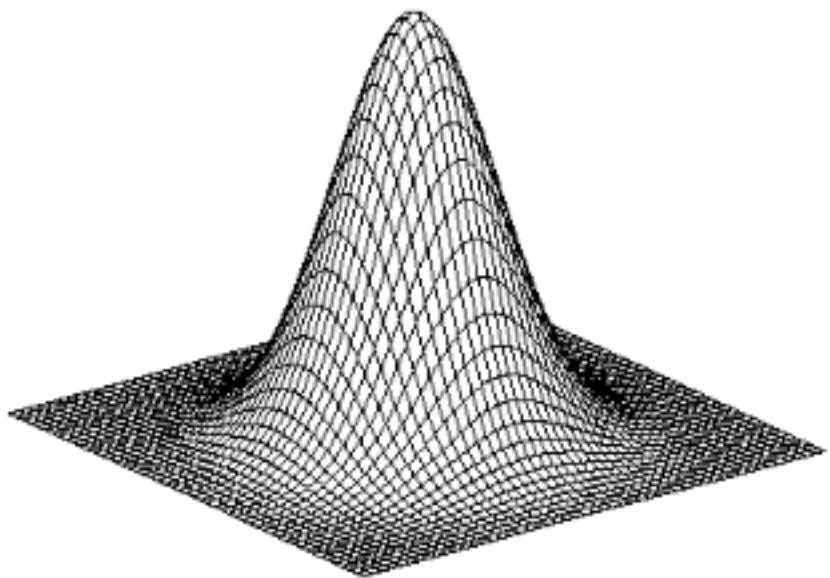
Derivatives of Filters

- Can optionally combine the blurring and differentiation steps using the theorem:

$$\frac{d}{dx}(f * g) = f * \frac{d}{dx}g$$

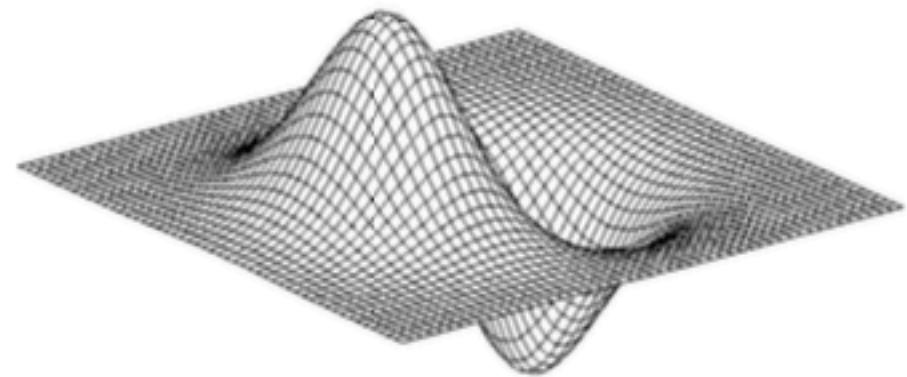
- **Algorithm 2 for simple edge detector:**
 1. Convolve with x derivative of Gaussian, gives E_x
 2. Convolve with y derivative of Gaussian, gives E_y
 3. Find gradient magnitude: $E = \| E_x^2 + E_y^2 \|$

Derivatives of Gaussian Filter



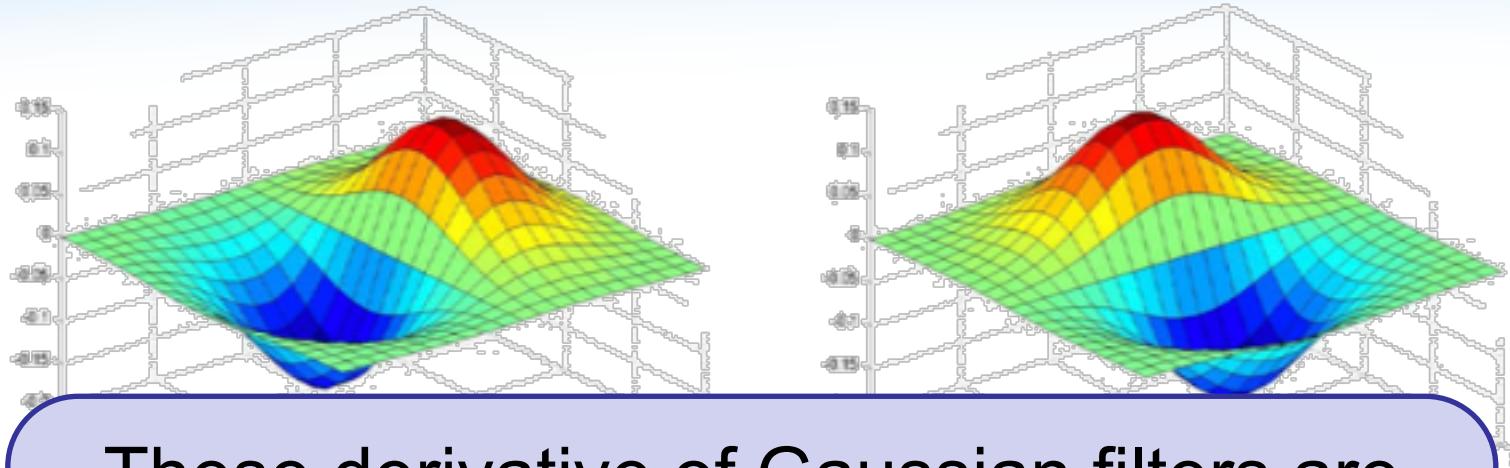
2D-gaussian

$$\ast [1 \ -1] =$$



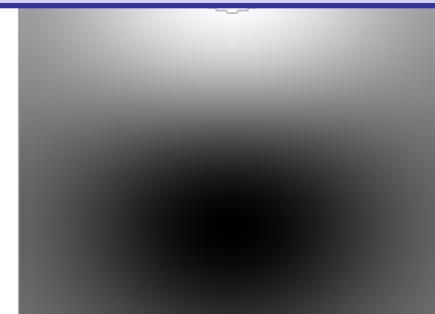
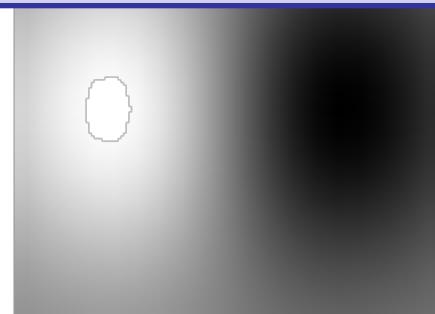
x - derivative

Derivatives of Gaussian Filter

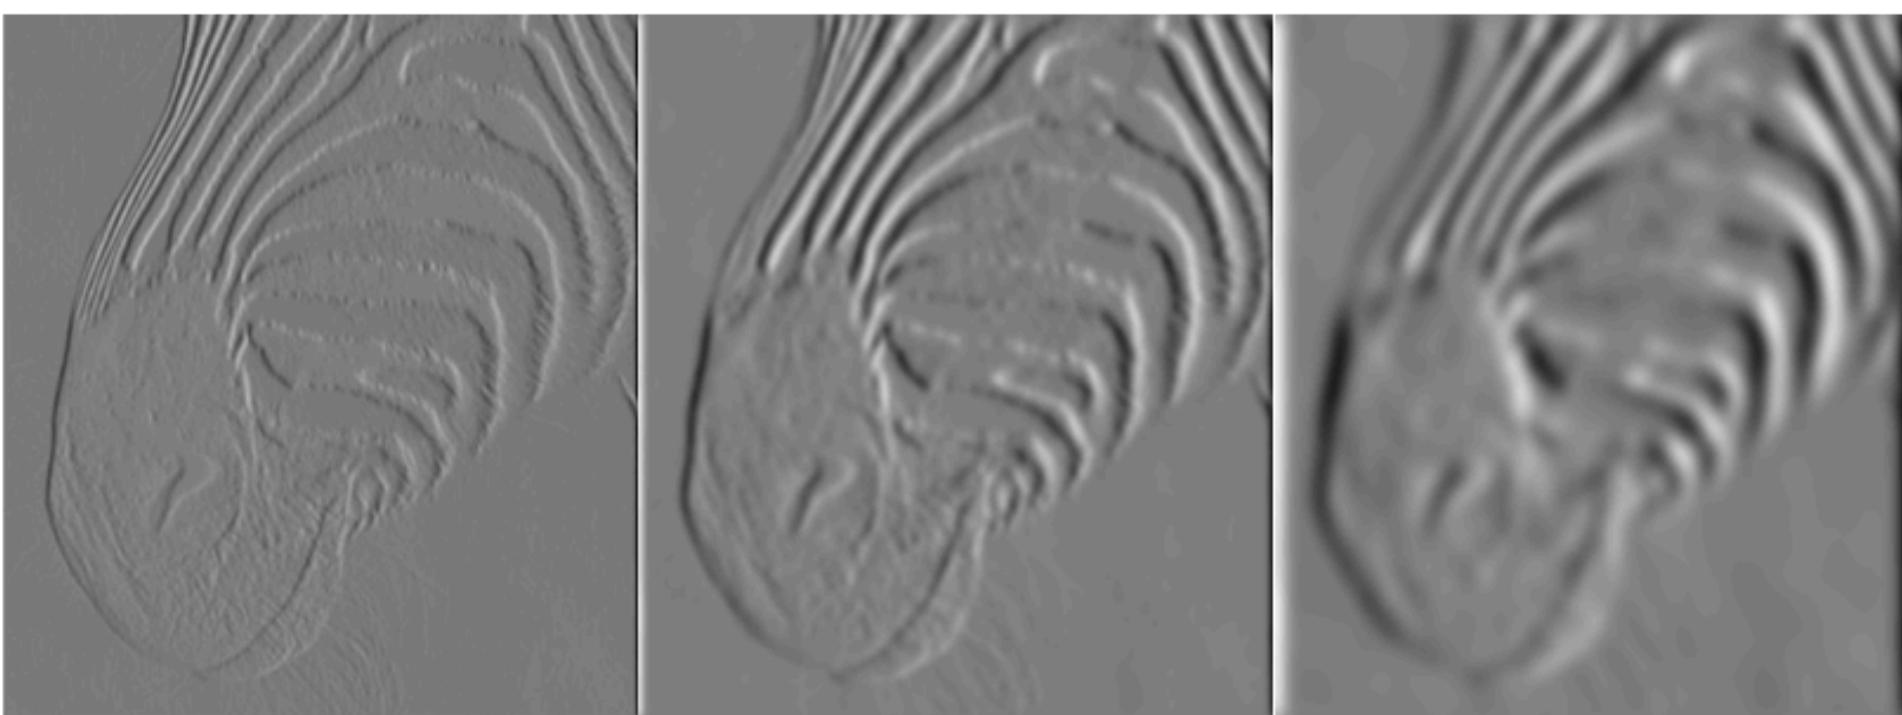


These derivative of Gaussian filters are separable, just like the Gaussian.

How does that help?



Effect of Gaussian Filter Width (σ)



1 pixel

3 pixels

7 pixels

- Smoothed derivative removes noise, but blurs edge. Also finds edges at different “scales”.

Remaining Issues



- The gradient magnitude is large along a thick “trail” or “ridge,” so how do we identify the actual edge points?
- How do we link the edge points to form curves?

Source: D. Forsyth

Outline

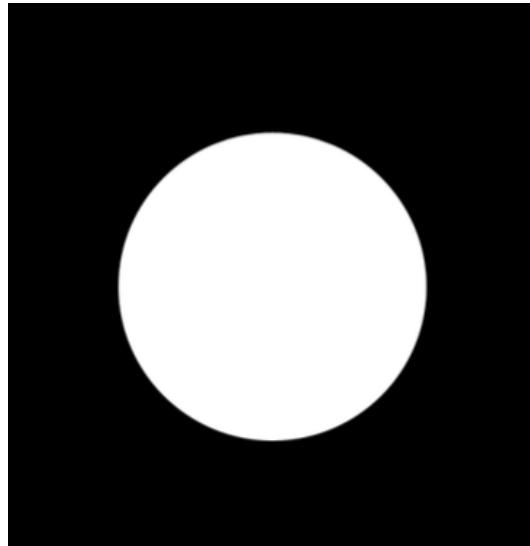
- Simple image processing
 - Greyscale
 - Brightness
 - Mirroring or “flipping”
- Filtering
 - Linear filters: cross-correlation and convolution
 - Gaussian filters
- Edge detection
 - Simple edge detector
 - **Canny edge detector**

Canny Edge Detector (in Project 1)

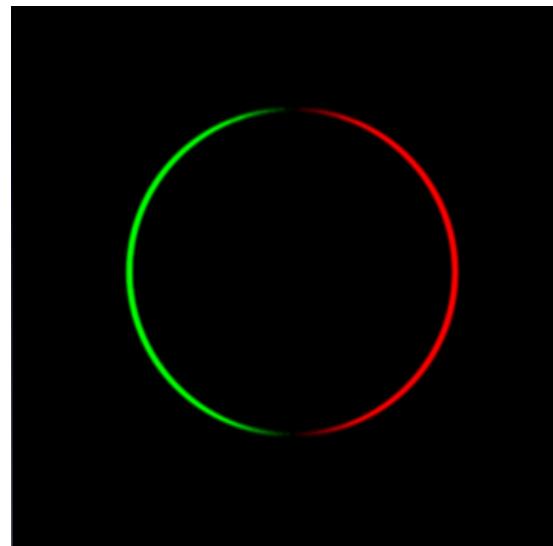
1. Smooth
2. Compute derivative
3. Non-maximum suppression
4. Thresholding

Canny Edge Detector

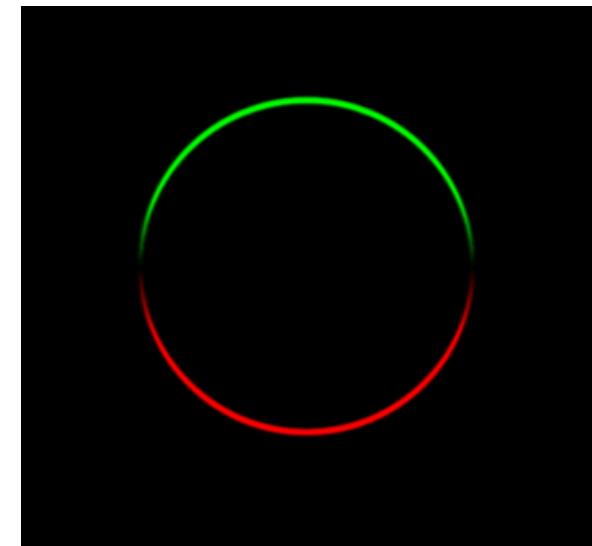
- First, smooth with a Gaussian with filter width σ
- Then compute x and y derivatives
- As we mentioned before the above 2 steps can be combined (using two derivative of Gaussian filters)



Input image



Smoothed x derivative



Smoothed y derivative

Canny Edge Detector

- **Non-maximum suppression:**
 - Eliminate all but local maxima in magnitude of gradient
 - At each pixel look along direction of gradient: if either neighbor is bigger, set to zero
 - In practice, quantize gradient directions to vertical, horizontal, two diagonals
 - Result: “thinned edge image.”

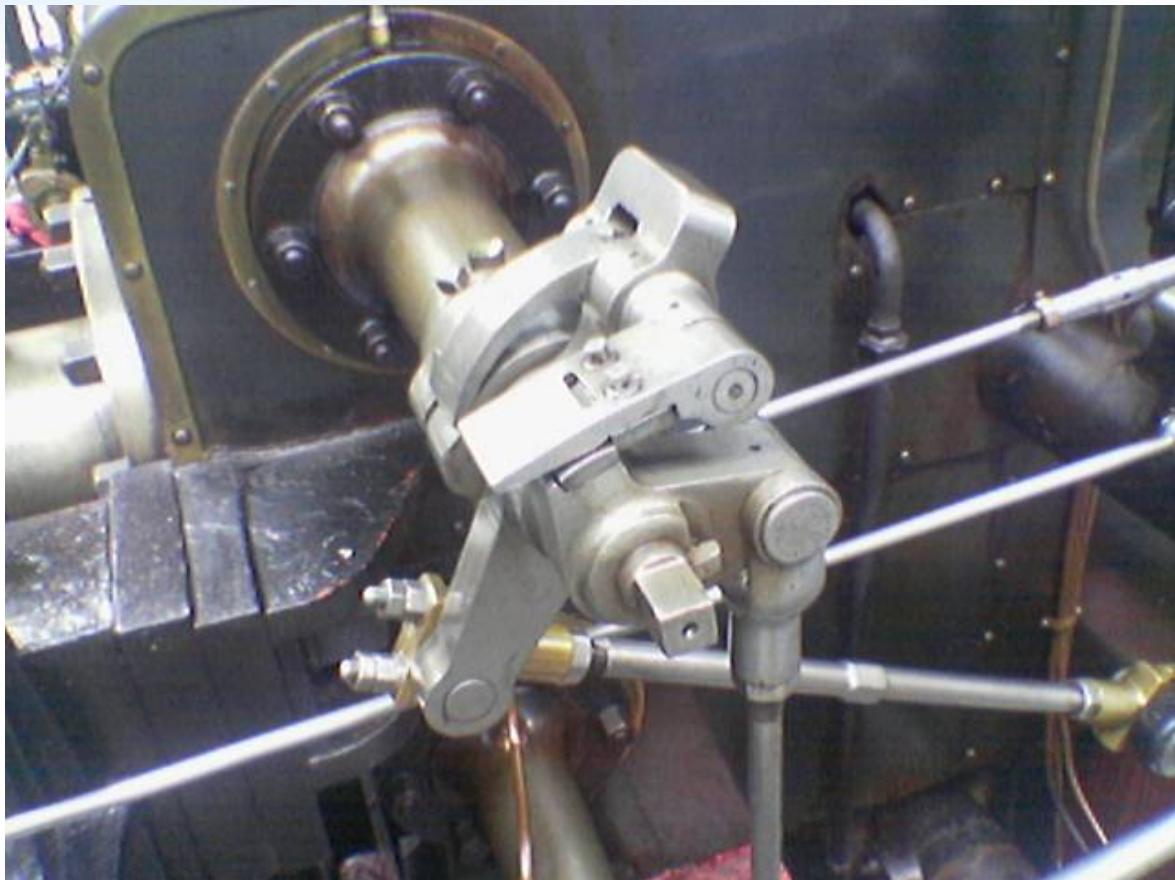
Canny Edge Detector

- Final stage: **thresholding**.
- Simplest: use a single threshold
- Better: use two thresholds
 - Mark pixels as “definitely not edge” if less than τ_{low} .
 - Mark pixels as “strong edge” if greater than τ_{high} .
 - Mark pixels as “weak edge” if within $[\tau_{low}, \tau_{high}]$.
 - Strong pixels are definitely part of the edge.
 - Weak pixels are debatable

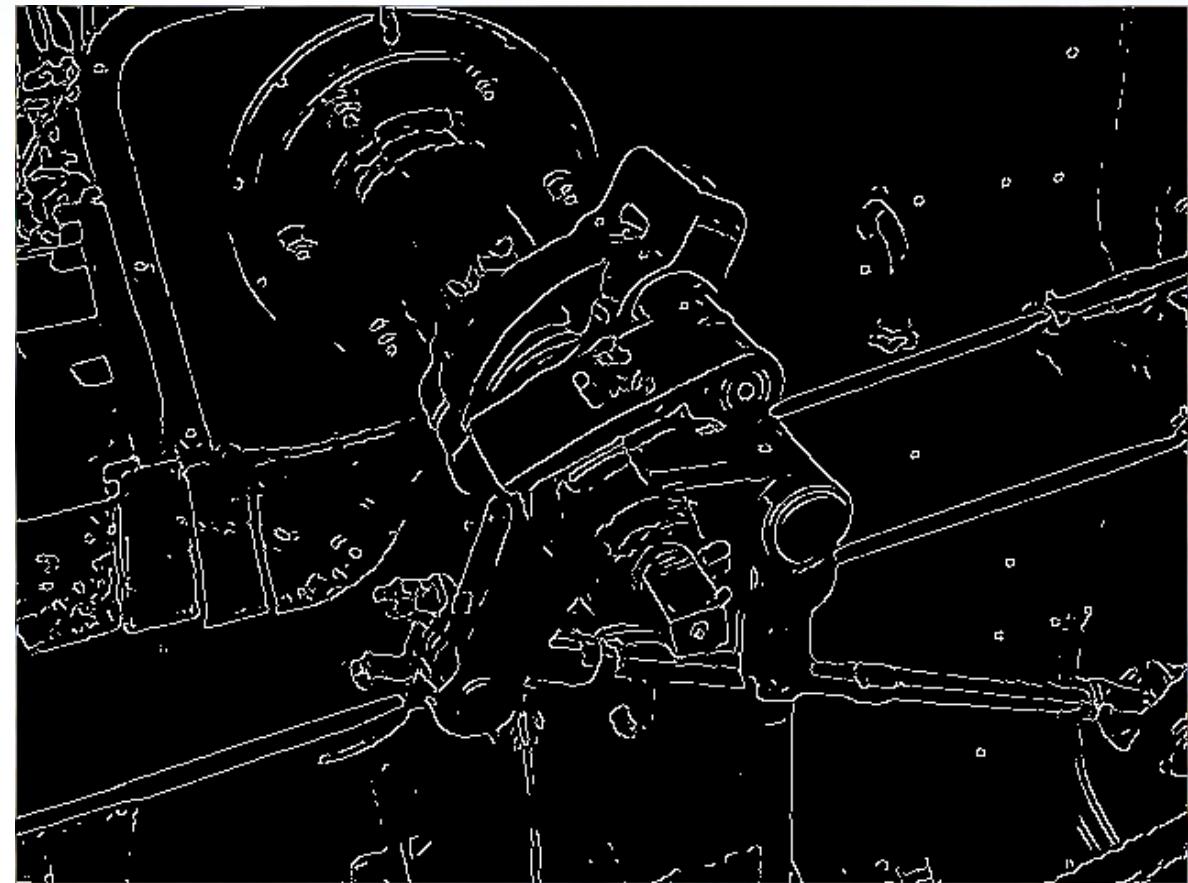
Canny Edge Detector

- Only include weak pixels connected in a chain to some strong pixel.
- How to do this?
 - Visit pixels in chains starting from the strong pixels. For each strong pixel, recursively visit the weak pixels that are in the 8 connected neighborhood around the strong pixel, and label those also as strong (and as edge).
 - Label as “not edge” any weak pixels that are not visited by this process.

Canny Edge Detector



Input image



Canny Edge Detector

From [Wikipedia](#)